



# CMPE 258, Deep Learning

## Logistic Regression

Feb 8, 2018

DMH 149A

**Taehee Jeong**


**Ph.D., Data Scientist**

# Recap

- Overfitting of polynomial regression
- Regularization
  - sum of square value, L2 norm, Ridge
  - sum of absolute value, L1 norm, Lasso
- Bias & Variance trade off

# Overfitting of polynomial regression

$$\hat{y} = W_0 + W_1x_1 + W_2x_1^2 + W_3x_1^3 + \dots$$



<Machine Learning, Emily Fox & Carlos Guestrin>

# Symptom of overfitting

- Very large value of regression coefficients,  $W$
- Lots of input features
- Small number of observations

# Avoiding Overfitting through Regularization

- **Sum of squares**

Ridge : L2 norm

$$W_0^2 + W_1^2 + W_2^2 + \dots = \sum_{j=1}^n W_j^2 = \|W\|_2^2$$

Cost function  $+\lambda\|W\|_2^2$

- **Sum of absolute value**

Lasso: L1 norm

$$|w_0| + |w_1| + |w_2| + \dots = \sum_{j=1}^n |w_j| = \|W\|_1$$

Cost function  $+\lambda\|W\|_1$

# Regularized linear regression

## Ridge regression (L2 penalty)

Cost function

$$J = \frac{1}{m} \sum_{i=1}^m (\hat{y}^i - y^i)^2 + \frac{\lambda}{m} \sum_{j=1}^n W_j^2$$

# Regularized linear regression

## Ridge regression (L2 penalty)

### Gradient

$$\frac{\partial J}{\partial W_0} = \frac{2}{m} \sum_{i=1}^m (\hat{y}^i - y^i) x_0^i$$

when  $j=0$

$$\frac{\partial J}{\partial W_j} = \frac{2}{m} \sum_{i=1}^m (\hat{y}^i - y^i) x_j^i + \frac{2\lambda}{m} W_j$$

when  $j \geq 1$

# Regularized linear regression

## Ridge regression (L2 penalty)

Gradient descent

Repeat {

$$W_0 = W_0 - \alpha \frac{2}{m} \sum_{i=1}^m (\hat{y}^i - y^i) x_0^i \quad \text{when } j=0$$

$$W_j = W_j - \alpha \left[ \frac{2}{m} \sum_{i=1}^m (\hat{y}^i - y^i) x_j^i + \frac{2\lambda}{m} W_j \right] \quad \text{when } j \geq 1$$

}



# Regularized linear regression

## Ridge regression (L2 penalty)

Matrix form

Cost function

$$J = \frac{1}{m} [(W \cdot X - Y)^T (W \cdot X - Y) + \lambda W^T \cdot W]$$

Gradient Descent

$$\frac{\partial J}{\partial W} = \frac{2}{m} [(X \cdot W - Y)^T \cdot X + \lambda W]$$

I: identity matrix

$$\frac{\partial J}{\partial W} = \frac{2}{m} [(X \cdot W - Y)^T \cdot X + \lambda I \cdot W]$$

# Regularized linear regression

## Ridge regression (L2 penalty)

Normalization equation for ridge regression


$$W = (X^T \cdot X + \lambda I)^{-1} \cdot X^T \cdot Y$$

when  $\lambda=0$ ,  $W = (X^T \cdot X)^{-1} \cdot X^T \cdot Y$

When  $\lambda = \text{infinite}$ ,  $W=0$

# Regularized linear regression

## Ridge regression (L2 penalty)



<Machine Learning, Emily Fox & Carlos Guestrin>

# Regularized linear regression

## Lasso regression (L1 penalty)


Matrix form

Cost function

$$J = \frac{1}{m} [(W \cdot X - Y)^T (W \cdot X - Y) + \lambda |W|]$$

Gradient Descent


Coordinate descent method



<Machine Learning, Emily Fox & Carlos Guestrin>

# Regularized linear regression

## Lasso regression (L1 penalty)



<Machine Learning, Emily Fox & Carlos Guestrin>

# Bias / Variance Trade off

Large  $\lambda$ : high bias, low variance

Small  $\lambda$ : low bias, high variance

## *Bias*


This part of the generalization error is due to wrong assumptions, such as assuming that the data is linear when it is actually quadratic. A high-bias model is most likely to underfit the training data.

## *Variance*

This part is due to the model's excessive sensitivity to small variations in the training data. A model with many degrees of freedom (such as a high-degree polynomial model) is likely to have high variance, and thus to overfit the training data.

<Hands-On ML, Aurelien Geron>

# Overfitting and regularization



<Machine Learning, Emily Fox & Carlos Guestrin>

# Assignment\_2: Any question?

Due 2/15

- 1 (30pts). Polynomial regression / overfitting / regularization
- 2 (30pts). Polynomial regression with train/validation/test
- 3 (40pts). Regularization with Tensorflow
  - using L2 penalty (Ridge)
  - using L1 penalty (Lasso)
  - using matrix (gradient descent)
  - using scikit-learn linear regression model
  - using TensorFlow gradient descent method



# Logistic Regression

## Binary classification

- Logistic regression is commonly used for binary classification.
- The output of classification is categorical value instead of continuous value.
- The output of binary classification is 1 or 0.
- For example,
  - Email: Spam / Not spam,
  - Online transaction: Fraudulent (Yes / No)

# Logistic Regression


Given  $x$ , the probability if  $y = 1$

$$\hat{y} = P(y = 1|x)$$

$$P = \sigma(W^T x + b)$$

Logistic function (sigmoid)

$$\sigma(z) = \frac{1}{1 + \exp(-z)}$$



<Machine Learning, Emily Fox & Carlos Guestrin>


# Logistic regression

## Model Prediction

$$P(y = 1|x)$$

$$= \frac{1}{1 + \exp(-W^T x - b)}$$

$$\hat{y} = \begin{cases} 1 & \text{if } P \geq 0.5 \\ 0 & \text{if } P < 0.5 \end{cases}$$



# Logistic Regression

## Cost function

$$J = -\frac{1}{m} \sum_{i=1}^m [y^i \log(\hat{y}^i) + (1 - y^i) \log(1 - \hat{y}^i)]$$

$$\text{If } y = 1 \quad J = -\frac{1}{m} \sum_{i=1}^m [y^i \log(\hat{y}^i)]$$

$$\text{If } y = 0 \quad J = -\frac{1}{m} \sum_{i=1}^m [\log(1 - \hat{y}^i)]$$

<Machine Learning, Emily Fox & Carlos Guestrin>

# Logistic Regression

Gradient (or derivative of cost function)

$$\frac{\partial J}{\partial W} = -\frac{1}{m} \sum_{i=1}^m [(\hat{y}^i - y^i) x_j^i]$$

$$\frac{\partial J}{\partial W} = -\frac{1}{m} \sum_{i=1}^m [(P - y^i) x_j^i]$$

$$\frac{\partial J}{\partial W} = -\frac{1}{m} \sum_{i=1}^m [(\sigma(W^T x + b) - y^i) x_j^i]$$

# Logistic Regression

## Regularization

Cost function

$$J = -\frac{1}{m} \sum_{i=1}^m [y^i \log(\hat{y}^i) + (1 - y^i) \log(1 - \hat{y}^i)] + \frac{2\lambda}{m} \sum_{j=1}^n W_j^2$$

gradient

$$\frac{\partial J}{\partial W_0} = -\frac{1}{m} \sum_{i=1}^m [(\hat{y}^i - y^i) x_j^i] \quad \text{for } j=0$$

$$\frac{\partial J}{\partial W_j} = -\frac{1}{m} \sum_{i=1}^m [(\hat{y}^i - y^i) x_j^i] + \frac{\lambda}{m} W_j \quad \text{for } j \geq 1$$

# Confusion matrix

		Predicted label	
		(+1)	(-1)
True label	(+1)	1503 <b>tp</b>	203 <b>fn</b>
	(-1)	905 <b>fp</b>	1845 <b>tn</b>

Accuracy :  $(tp + tn) / (tp + fp + fn + tn) * 100$

Precision :  $tp / \text{predict positive} = tp / (tp + fp) * 100$

Recall (true positive rate) :  $tp / \text{actual positive} = tp / (tp + fn) * 100$

True negative rate :  $tn / \text{actual negative} = tn / (fp + tn) * 100$

# Precision

Fraction of positive predictions that are correct.

$$\text{precision} = \frac{\text{\# true positives}}{\text{\# true positives} + \text{\# false positives}}$$




# Recall

Fraction of positive data points correctly classified

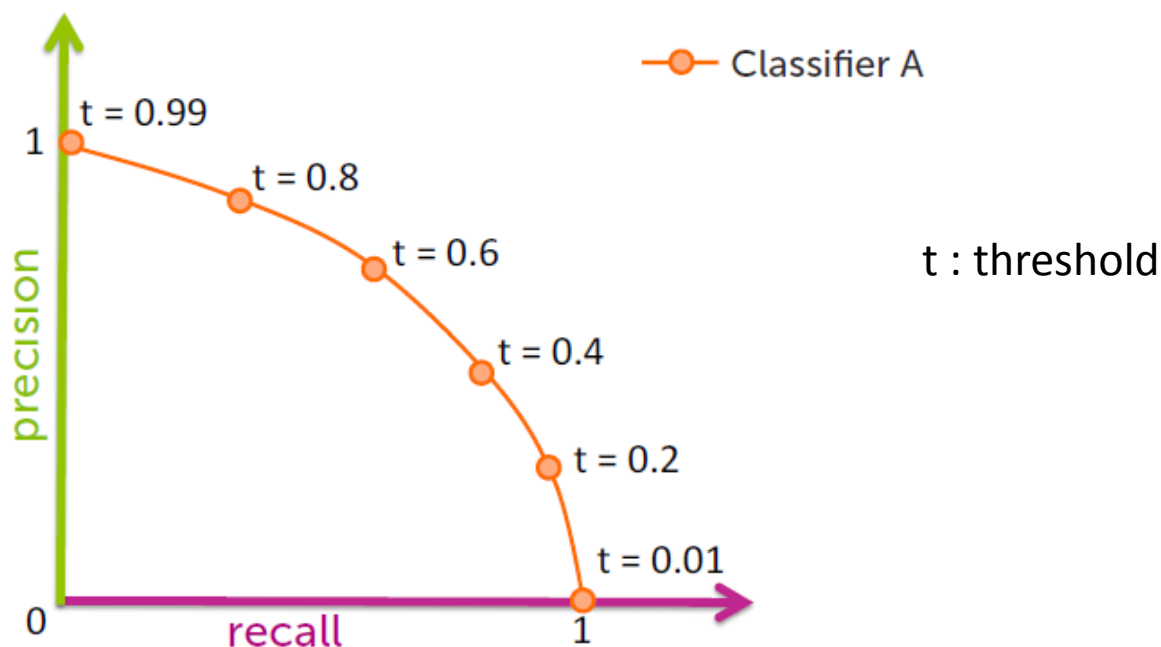
$$\text{Recall} = \frac{\# \text{ true positives}}{\# \text{ true positives} + \# \text{ false negatives}}$$

# Tradeoff Precision & Recall



<Machine Learning, Emily Fox & Carlos Guestrin>

# Precision & Recall curve



<Machine Learning, Emily Fox & Carlos Guestrin>

# Mini-batch Gradient Descent

- Batch Gradient Descent

Computes the gradients based on full training set

Ex) Offline learning

- Stochastic Gradient Descent

Computes just one instance

Ex) Online learning

- Mini-batch Gradient Descent


Computes the gradients on small random sets of instances

<Hands-On ML, Aurelien Geron>

# Batch vs online learning

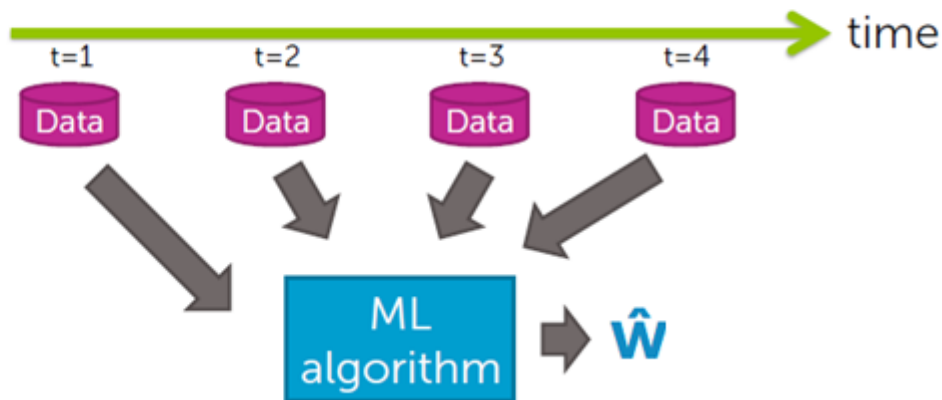
## Batch learning

- All data is available at start of training time



## Online learning

- Data arrives (streams in) over time
  - Must train model as data arrives!



<Machine Learning, Emily Fox & Carlos Guestrin>

# Stochastic gradient descent

**Batch  
Gradient  
descent**

$$\frac{\partial J}{\partial W} = -\frac{1}{m} \sum_{i=1}^m [(\hat{y}^i - y^i)x_j^i] \quad (\text{Sum over data points})$$

**Stochastic  
Gradient  
descent**

$$\frac{\partial J}{\partial W} = -(\hat{y}^i - y^i)x_j^i \quad (\text{Each time, pick different data point})$$

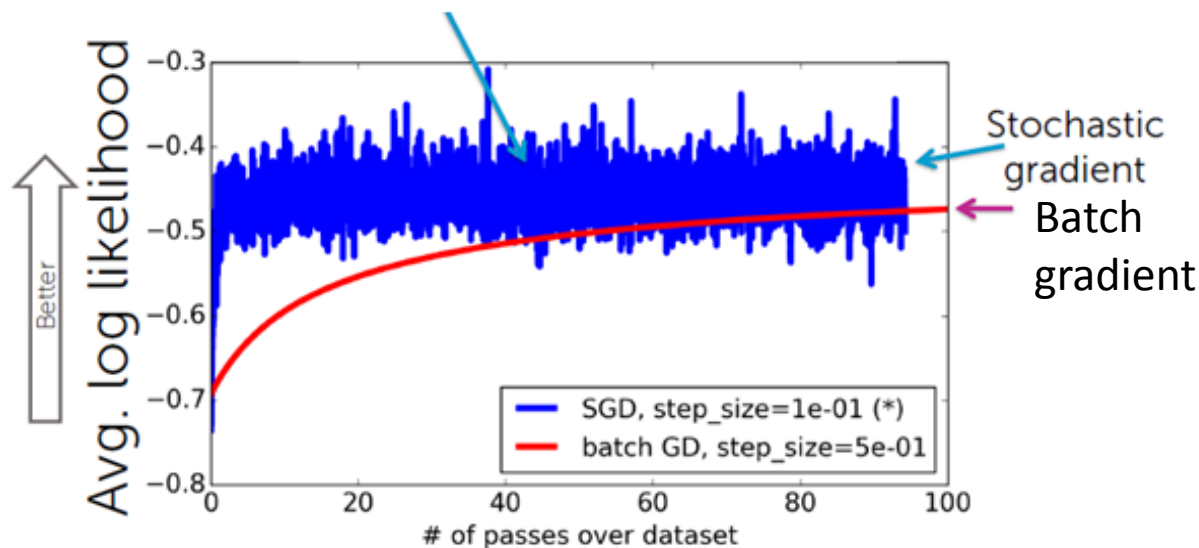
# Batch vs. Stochastic gradient

Algorithm	Time per iteration	Total time to convergence for large data		Sensitivity to parameters
		In theory	In practice	
Batch Gradient	Slow for large data	Slower	Often slower	Moderate
Stochastic gradient	Always fast	Faster	Often faster	Very high

<Machine Learning, Emily Fox & Carlos Guestrin>

# Batch gradient descent vs stochastic gradient descent

Consider only 'average' value in stochastic gradient



<Machine Learning, Emily Fox & Carlos Guestrin>



# Stochastic gradient

Tiny change to gradient ascent

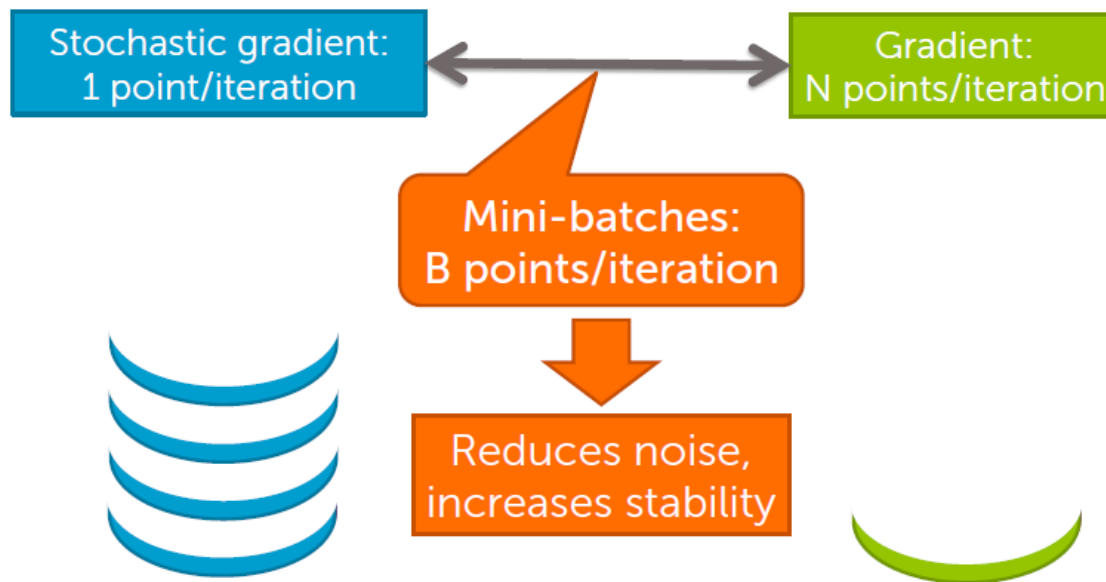
Much better scalability

Huge impact in real-world

Very tricky to get right in practice

<Machine Learning, Emily Fox & Carlos Guestrin>

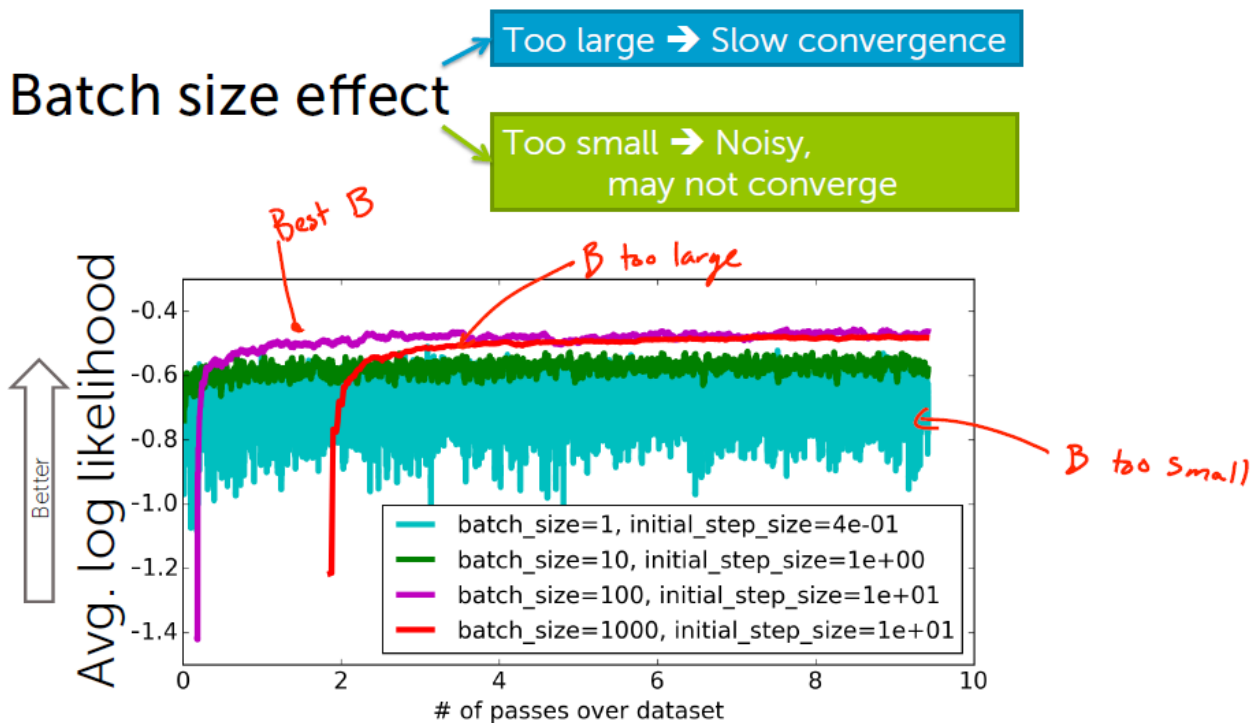
# Batch / stochastic: two extremes



<Machine Learning, Emily Fox & Carlos Guestrin>

# Batch size effect in mini-batch

## Batch size effect



<Machine Learning, Emily Fox & Carlos Guestrin>

# Regression with Tensorflow

Load data and set up X, y variable

```
import tensorflow as tf
import numpy as np
from sklearn.datasets import fetch_california_housing
housing = fetch_california_housing()
m,n = housing.data.shape
mean = np.mean(housing.data,axis=0)
std = np.std(housing.data, axis=0)
normal_housing_data = (housing.data - mean)/std
normal_housing_data_plus_bias = np.c_[np.ones((m,1)),normal_housing_data]
X = tf.constant(normal_housing_data_plus_bias, dtype = tf.float32, name = "x")
y = tf.constant(housing.target.reshape(-1,1), dtype = tf.float32, name = "y")
```

# Regression with Tensorflow

## L2 regularization

```
theta = tf.Variable(tf.random_uniform([n+1,1],-1.0,1.0), name = "theta")
y_pred = tf.matmul(X,theta,name = "Predictions")
error = y_pred - y
rmse = tf.sqrt(tf.reduce_mean(tf.square(error)), name = "rmse")
scale = 0.1
learning_rate = 0.01
base_loss = tf.reduce_mean(tf.square(error), name = "loss")
reg_loss = tf.reduce_sum(tf.square(theta))
loss = tf.add(base_loss, scale/m*reg_loss)
gradients = 2/m * tf.add(tf.matmul(tf.transpose(X),error),scale*theta)
training_op = tf.assign(theta, theta - learning_rate * gradients)
```

# Regression with Tensorflow

## L1 regularization

```
theta = tf.Variable(tf.random_uniform([n+1,1],-1.0,1.0), name = "theta")
y_pred = tf.matmul(X,theta,name = "Predictions")
error = y_pred - y
rmse = tf.sqrt(tf.reduce_mean(tf.square(error)), name = "rmse")
scale = 0.1
base_loss = tf.reduce_mean(tf.square(error), name = "loss")
reg_loss = tf.reduce_sum(tf.abs(theta))
loss = tf.add(base_loss, scale/m*reg_loss)
gradients = 2/m * tf.add(tf.matmul(tf.transpose(X),error),scale*theta)
training_op = tf.assign(theta, theta - learning_rate * gradients)
```

# Summary

- Logistic regression
- Binary classification
- Confusion matrix: Accuracy, Precision, Recall
- Mini-batch gradient descent
- Regression with tensorflow