



CMPE 258, Deep Learning

Convolutional layer

March 20, 2018

DMH 149A

Taehee Jeong

Ph.D., Data Scientist

Assignment_4

Deep Neural Network with tensorflow

Due day is Sunday, March 18th.

There is penalty for late submission or re-submission after due day.

In Assignment_3 and mid-term exam, we used pandas and numpy.

For Assignment_4, we will use Tensorflow.

Please build Deep Neural Network with two hidden layers.

For activation function, please use relu or elu for hidden layers, sigmoid for output layer.

For weight initialization, please use xavier initialization.

For optimization, please use adam optimization.

For regularization, please use dropout.


As the final output, please plot train accuracy and test accuracy with probability (0.1 ~ 0.9) of dropout.

Today's lesson

- Compare CNN vs. DNN
- 2D image filter
- Convolution calculation
- Convolution on RGB images
- Multiple filters
- Size of matrix in convolution layers

Digital representation of an image

- Grayscale image is a matrix of pixels (**picture elements**)
- Dimensions of this matrix are called image resolution (e.g. 20 x 20)
- Each pixel stores its brightness (or **intensity**) ranging from 0 to 255, 0 intensity corresponds to black color
- Color images store pixel intensities for 3 channels: red, green and blue




| | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | ... | 391 | 392 | 393 | 394 | 395 | 396 | 397 | 398 | 399 | y |
|---|---|---|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|---|
| 0 | 0 | 0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | ... | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0 | 5 |
| 1 | 0 | 0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | ... | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0 | 9 |
| 2 | 0 | 0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | ... | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0 | 7 |
| 3 | 0 | 0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | ... | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0 | 6 |
| 4 | 0 | 0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | ... | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0 | 5 |

<Intro to deep learning, national research university>


Normalize input pixels

$$x_{norm} = \frac{x}{255} - 0.5$$

| | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | ... | 391 | 392 | 393 | 394 | 395 | 396 | 397 | 398 | 399 | y |
|---|---|---|---|---|---|---|---|---|---|---|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | ... | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0 5 |
| 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | ... | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0 9 |
| 2 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | ... | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0 7 |
| 3 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | ... | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0 6 |
| 4 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | ... | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0 5 |



Why not Deep Neural Network?




Why not Deep Neural Network?

CNNs solve this problem using **partially connected layers**.



$$1000 \times 1000 \times 3 = 3 \text{ million}$$

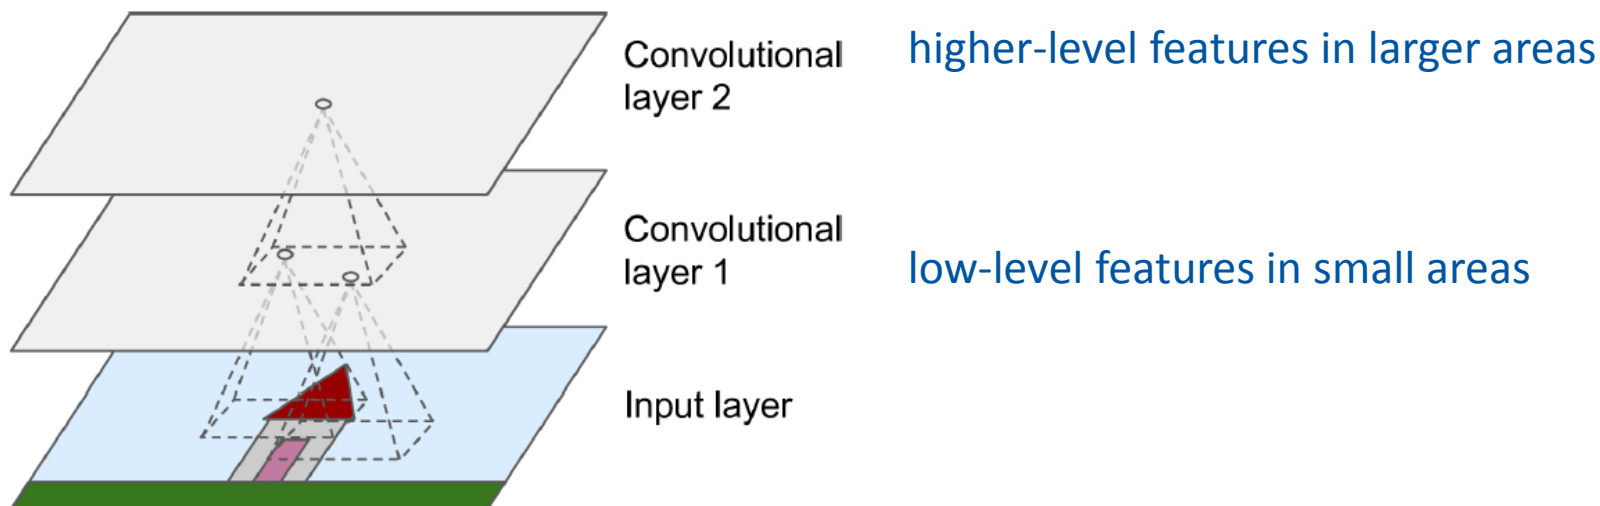


<Deep Neural Network, Andrew Ng>

Convolutional layer

Inspired by the architecture of visual cortex by David H. Hubel and Torsten wiesel


Hierarchical structure




<Hands-On ML, A. Geron>

Convolutional layer

low-level features



high-level features




higher-level features



<Deep Neural Network, Andrew Ng>

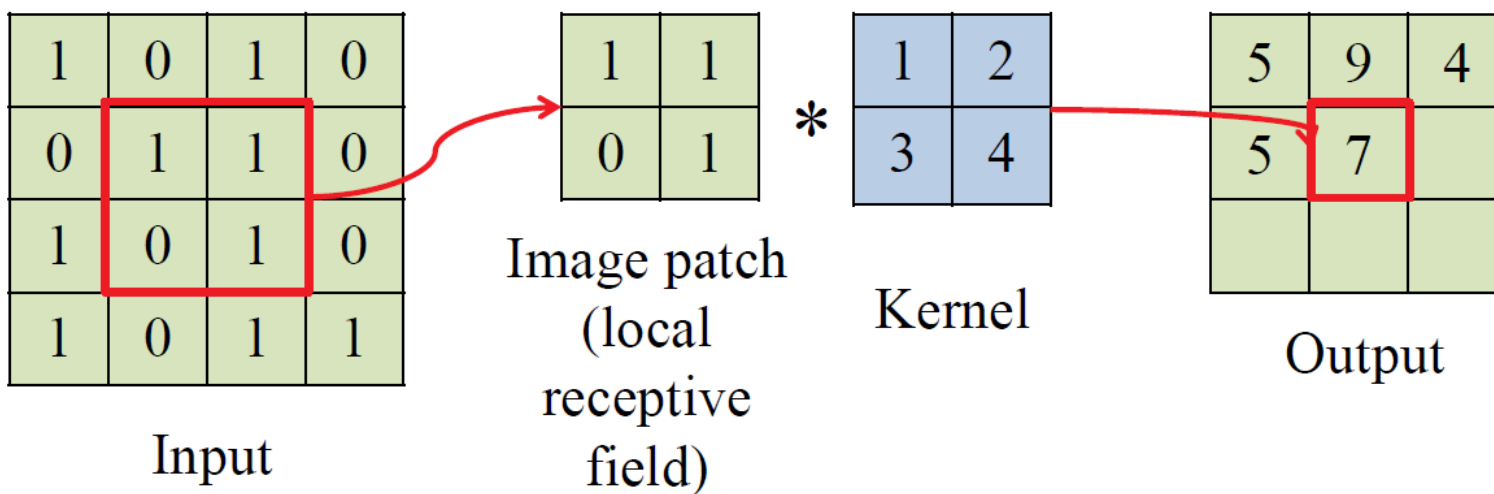
Convolution

Convolution is a dot product of a **kernel** (or filter) and a patch of an image (**local receptive field**) of the same size



<Intro to deep learning, national research university>

Convolution



<Intro to deep learning, national research university>

Vertical edge detection

| | | | | | |
|----|-----------|-----------|----------|---|---|
| 10 | 10 | 10 | 0 | 0 | 0 |
| 10 | 10 | 10 | 0 | 0 | 0 |
| 10 | 10 | 10 | 0 | 0 | 0 |
| 10 | <u>10</u> | <u>10</u> | <u>0</u> | 0 | 0 |
| 10 | <u>10</u> | <u>10</u> | <u>0</u> | 0 | 0 |
| 10 | <u>10</u> | <u>10</u> | <u>0</u> | 0 | 0 |

6x6

3x3

| | | |
|---|---|----|
| 1 | 0 | -1 |
| 1 | 0 | -1 |
| 1 | 0 | -1 |

3x3

*


| | | | |
|---|----|----|---|
| 0 | 30 | 30 | 0 |
| 0 | 30 | 30 | 0 |
| 0 | 30 | 30 | 0 |
| 0 | 30 | 30 | 0 |

4x4

<Deep Neural Network, Andrew Ng>


Vertical edge detection

| | | | | | |
|----|----|----|---|---|---|
| 10 | 10 | 10 | 0 | 0 | 0 |
| 10 | 10 | 10 | 0 | 0 | 0 |
| 10 | 10 | 10 | 0 | 0 | 0 |
| 10 | 10 | 10 | 0 | 0 | 0 |
| 10 | 10 | 10 | 0 | 0 | 0 |
| 10 | 10 | 10 | 0 | 0 | 0 |




*

| | | |
|---|---|----|
| 1 | 0 | -1 |
| 1 | 0 | -1 |
| 1 | 0 | -1 |




=

| | | | |
|---|----|----|---|
| 0 | 30 | 30 | 0 |
| 0 | 30 | 30 | 0 |
| 0 | 30 | 30 | 0 |
| 0 | 30 | 30 | 0 |




| | | | | | |
|---|---|---|----|----|----|
| 0 | 0 | 0 | 10 | 10 | 10 |
| 0 | 0 | 0 | 10 | 10 | 10 |
| 0 | 0 | 0 | 10 | 10 | 10 |
| 0 | 0 | 0 | 10 | 10 | 10 |
| 0 | 0 | 0 | 10 | 10 | 10 |
| 0 | 0 | 0 | 10 | 10 | 10 |




*

| | | |
|---|---|----|
| 1 | 0 | -1 |
| 1 | 0 | -1 |
| 1 | 0 | -1 |



=

| | | | |
|---|-----|-----|---|
| 0 | -30 | -30 | 0 |
| 0 | -30 | -30 | 0 |
| 0 | -30 | -30 | 0 |
| 0 | -30 | -30 | 0 |



<Deep Neural Network, Andrew Ng>

Horizontal edge detection

| | | | | | |
|----|----|----|----|----|----|
| 10 | 10 | 10 | 0 | 0 | 0 |
| 10 | 10 | 10 | 0 | 0 | 0 |
| 10 | 10 | 10 | 0 | 0 | 0 |
| 0 | 0 | 0 | 10 | 10 | 10 |
| 0 | 0 | 0 | 10 | 10 | 10 |
| 0 | 0 | 0 | 10 | 10 | 10 |


6x6

*

| | | |
|----|----|----|
| 1 | 1 | 1 |
| 0 | 0 | 0 |
| -1 | -1 | -1 |

=

| | | | |
|----|----|-----|-----|
| 0 | 0 | 0 | 0 |
| 30 | 10 | -10 | -30 |
| 30 | 10 | -10 | -30 |
| 0 | 0 | 0 | 0 |



<Deep Neural Network, Andrew Ng>

Edge detection

Kernel

$$\begin{matrix} * & \begin{array}{|c|c|c|} \hline -1 & -1 & -1 \\ \hline -1 & 8 & -1 \\ \hline -1 & -1 & -1 \\ \hline \end{array} & = \end{matrix}$$

Edge detection



Original
image

Sums up to 0 (black color)
when the patch is a solid fill



Sharpening

Original image

Kernel

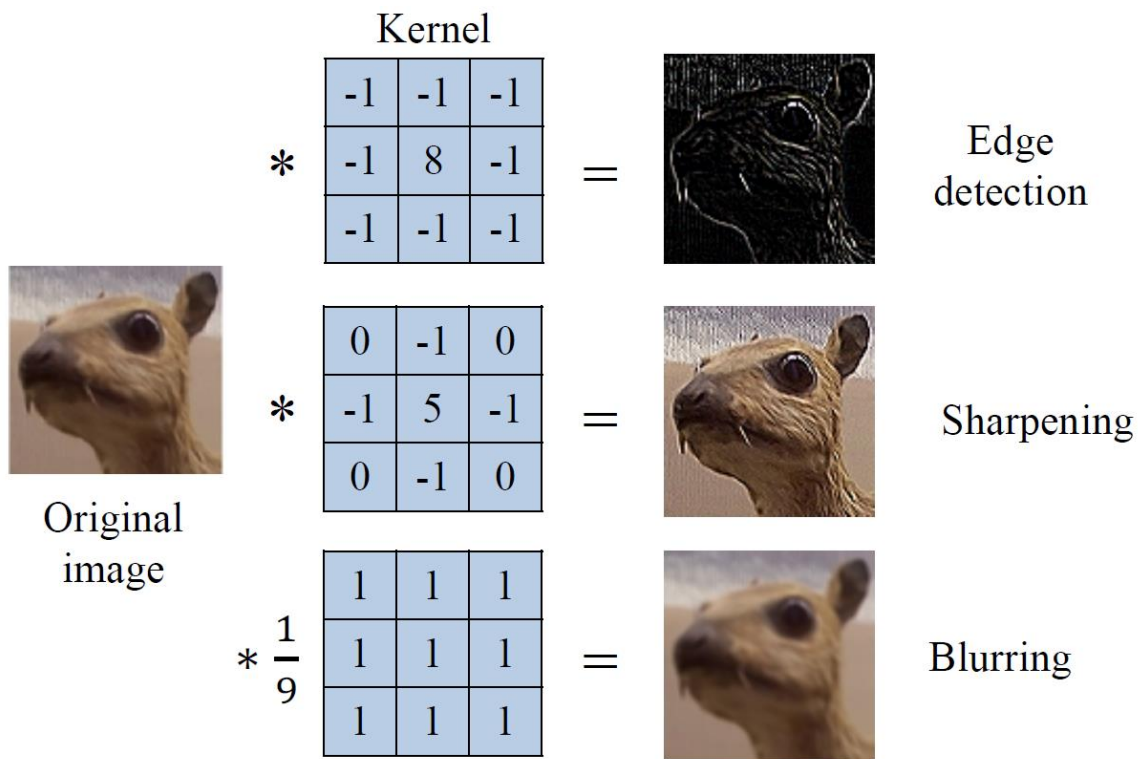
$\begin{matrix} * & \begin{matrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{matrix} & = & \begin{matrix} \text{Edge} \\ \text{detection} \end{matrix} \end{matrix}$

$\begin{matrix} * & \begin{matrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{matrix} & = & \begin{matrix} \text{Sharpening} \end{matrix} \end{matrix}$

Doesn't change an image for solid fills





Adds a little intensity on the edges

Blurring




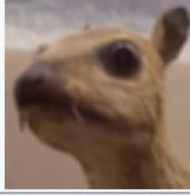

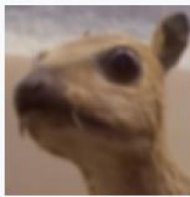
<Intro to deep learning, national research university>

Kernel (image processing)

| Operation | Kernel | Image result |
|----------------|---|--|
| Identity | $\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$ |  |
| Edge detection | $\begin{bmatrix} 1 & 0 & -1 \\ 0 & 0 & 0 \\ -1 & 0 & 1 \end{bmatrix}$ |  |
| | $\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$ |  |
| | $\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$ |  |

<Wikipedia>

Kernel (image processing)

| | | |
|---|--|--|
| Sharpen | $\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$ |  |
| Box blur (normalized) | $\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$ |  |
| Gaussian blur 3 × 3 (approximation) | $\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$ |  |
| Gaussian blur 5 × 5 (approximation) | $\frac{1}{256} \begin{bmatrix} 1 & 4 & 6 & 4 & 1 \\ 4 & 16 & 24 & 16 & 4 \\ 6 & 24 & 36 & 24 & 6 \\ 4 & 16 & 24 & 16 & 4 \\ 1 & 4 & 6 & 4 & 1 \end{bmatrix}$ |  |

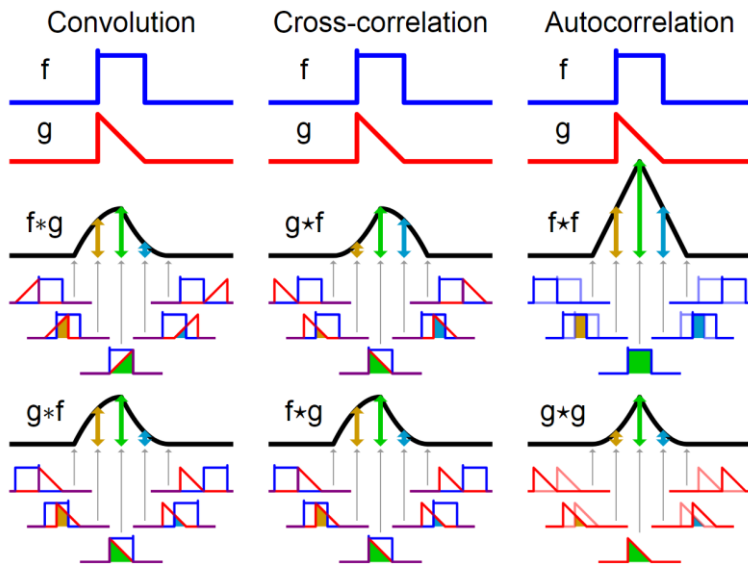
<Wikipedia>

Advantages of a CNN over a fully connected DNN for image classification

- CNN has fewer parameters than a fully connected DNN
- CNN can detect a particular feature anywhere on the image instead of detecting only in that particular location.
- CNN's architecture is able to identify pixels' organization (how each pixel is located).

Convolution

Convolutional layers actually use cross-correlation. <Hands-On ML, A. Geron>



<Wikipedia>

Convolution

$$y(n) = \sum_{k=-\infty}^{\infty} x(k)h(n-k)$$

<Digital signal processing, J.G., Proakis et al.>

Convolution is a mathematical operation that slides one function over another and measures the integral of their pointwise multiplication. It is heavily used in signal processing.

<Hands-On ML, A. Geron>

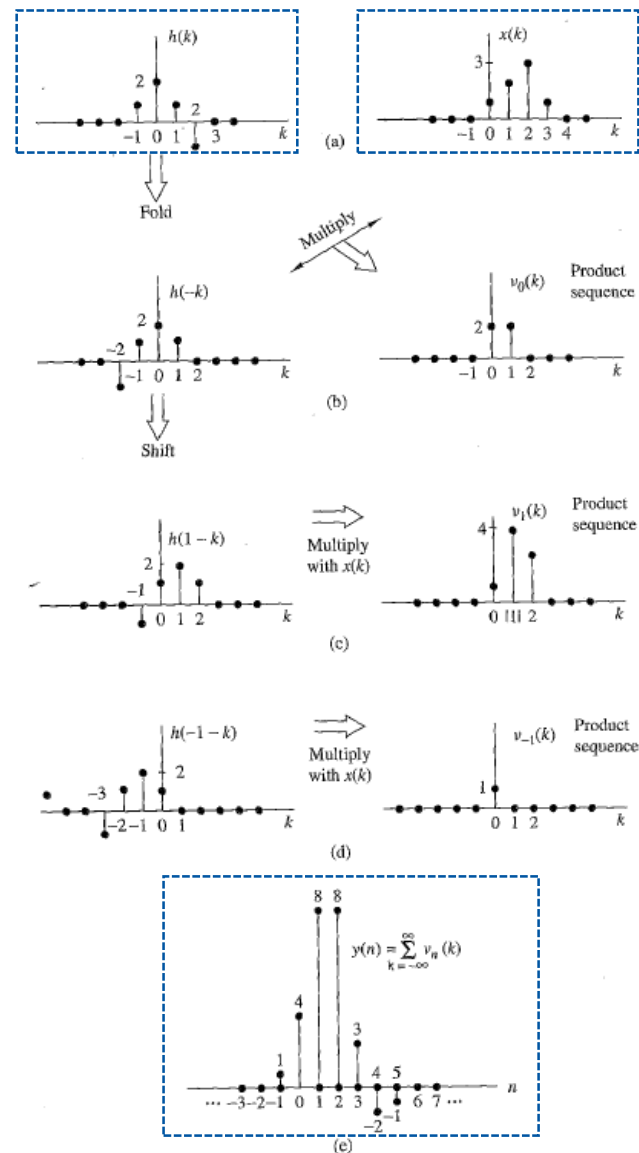



Image classification using Convolution Neural Network



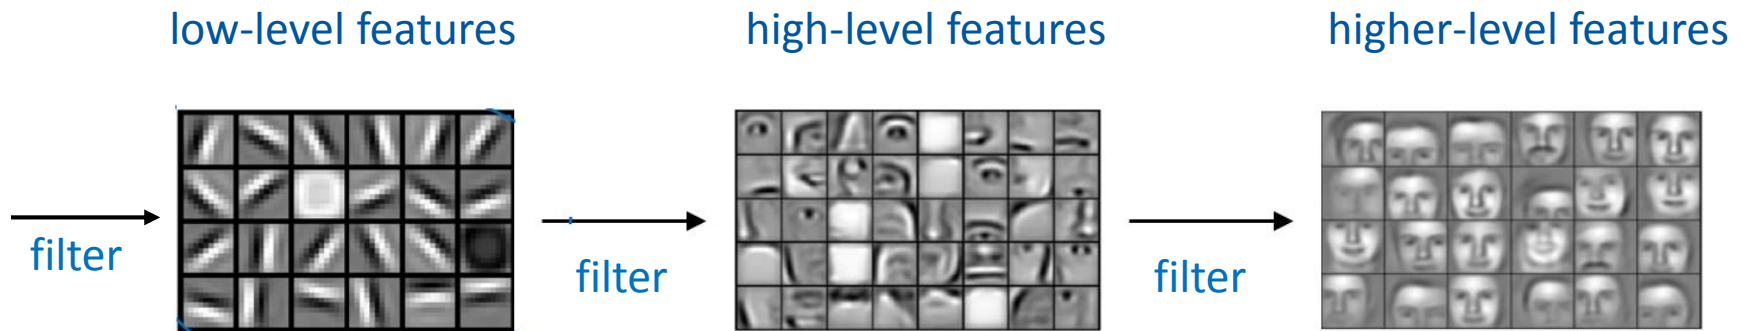
<Ralph Wittig, Power-Efficient Machine Learning using FPGAs on Power systems, OpenPOWERSummit>

Types of layer in a convolutional network

- Convolution
- Pooling
- Fully connected


Convolutional layer

What kind of filters should we use?



<Deep Neural Network, Andrew Ng>

Convolution



Input matrix

| | | | |
|----------|----------|----------|----------|
| X_1 | X_2 | X_3 | X_4 |
| X_5 | X_6 | X_7 | X_8 |
| X_9 | X_{10} | X_{11} | X_{12} |
| X_{13} | X_{14} | X_{15} | X_{16} |

Size: 4 x 4

convolution
*

filter

| | | |
|-------|-------|-------|
| W_1 | W_2 | W_3 |
| W_4 | W_5 | W_6 |
| W_7 | W_8 | W_9 |

Size: 3 x 3

Output matrix

What is the size?

Convolution

Input matrix

| | | | |
|----------|----------|----------|----------|
| X_1 | X_2 | X_3 | X_4 |
| X_5 | X_6 | X_7 | X_8 |
| X_9 | X_{10} | X_{11} | X_{12} |
| X_{13} | X_{14} | X_{15} | X_{16} |

Size: 4 x 4

convolution
*

filter

| | | |
|-------|-------|-------|
| W_1 | W_2 | W_3 |
| W_4 | W_5 | W_6 |
| W_7 | W_8 | W_9 |

Size: 3 x 3

Output matrix

=

| | |
|-------|-------|
| Z_1 | Z_2 |
| Z_3 | Z_4 |

Size: 2 x 2

Convolution

Input matrix

| | | | |
|----------|----------|----------|----------|
| X_1 | X_2 | X_3 | X_4 |
| X_5 | X_6 | X_7 | X_8 |
| X_9 | X_{10} | X_{11} | X_{12} |
| X_{13} | X_{14} | X_{15} | X_{16} |

Size: 4 x 4

convolution
*

filter

| | |
|-------|-------|
| W_1 | W_2 |
| W_3 | W_4 |

Size: 2 x 2

Output matrix

What is the size?

Convolution

Input matrix

| | | | |
|----------|----------|----------|----------|
| X_1 | X_2 | X_3 | X_4 |
| X_5 | X_6 | X_7 | X_8 |
| X_9 | X_{10} | X_{11} | X_{12} |
| X_{13} | X_{14} | X_{15} | X_{16} |

Size: 4 x 4

convolution
*

filter

| | |
|-------|-------|
| W_1 | W_2 |
| W_3 | W_4 |

Size: 2 x 2


Output matrix

=

Size: 3 x 3

| | | |
|-------|-------|-------|
| Z_1 | Z_2 | Z_3 |
| Z_4 | Z_5 | Z_6 |
| Z_7 | Z_8 | Z_9 |


Convolution calculation



<image Convolution>

Machinelearningguru.com/computer_vision/basics/convolution/image_convolution_1.html

Convolution calculation



<image Convolution>

Machinelearningguru.com/computer_vision/basics/convolution/image_convolution_1.html

Convolution calculation:pre-activation

| | | | |
|-----------------|-----------------|-----------------|-----------------|
| X ₁ | X ₂ | X ₃ | X ₄ |
| X ₅ | X ₆ | X ₇ | X ₈ |
| X ₉ | X ₁₀ | X ₁₁ | X ₁₂ |
| X ₁₃ | X ₁₄ | X ₁₅ | X ₁₆ |

 $*$

| | | |
|----------------|----------------|----------------|
| W ₁ | W ₂ | W ₃ |
| W ₄ | W ₅ | W ₆ |
| W ₇ | W ₈ | W ₉ |

 $=$

| | |
|----------------|----------------|
| Z ₁ | Z ₂ |
| Z ₃ | Z ₄ |

$$Z_1 = X_1 \times W_1 + X_2 \times W_2 + X_3 \times W_3 \\ + X_5 \times W_4 + X_6 \times W_5 + X_7 \times W_6 \\ + X_9 \times W_7 + X_{10} \times W_8 + X_{11} \times W_9$$

$$Z_3 = X_5 \times W_4 + X_6 \times W_5 + X_7 \times W_6 \\ + X_9 \times W_7 + X_{10} \times W_8 + X_{11} \times W_9 \\ + X_{13} \times W_{14} + X_{10} \times W_8 + X_{15} \times W_9$$

$$Z_2 = X_2 \times W_1 + X_3 \times W_2 + X_4 \times W_3 \\ + X_6 \times W_4 + X_7 \times W_5 + X_8 \times W_6 \\ + X_{10} \times W_7 + X_{11} \times W_8 + X_{12} \times W_9$$

$$Z_4 = X_6 \times W_4 + X_7 \times W_5 + X_8 \times W_6 \\ + X_{10} \times W_7 + X_{11} \times W_8 + X_{12} \times W_9 \\ + X_{14} \times W_7 + X_{15} \times W_8 + X_{16} \times W_9$$

Convolution: activation

Input matrix

| | | | |
|----------|----------|----------|----------|
| X_1 | X_2 | X_3 | X_4 |
| X_5 | X_6 | X_7 | X_8 |
| X_9 | X_{10} | X_{11} | X_{12} |
| X_{13} | X_{14} | X_{15} | X_{16} |

Size: 4 x 4

filter

| | | |
|-------|-------|-------|
| W_1 | W_2 | W_3 |
| W_4 | W_5 | W_6 |
| W_7 | W_8 | W_9 |

Size: 3 x 3

convolution
*

g : activation function

b : bias

=

| | |
|--------------|--------------|
| $g(Z_1 + b)$ | $g(Z_2 + b)$ |
| $g(Z_3 + b)$ | $g(Z_4 + b)$ |


=

Output matrix

| | |
|-------|-------|
| a_1 | a_2 |
| a_3 | a_4 |

Size: 2 x 2

Edge pixels?



What about information from edge pixels?

Padding

| | | | | | |
|---|----------|----------|----------|----------|---|
| 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | x_1 | x_2 | x_3 | x_4 | 0 |
| 0 | x_5 | x_6 | x_7 | x_8 | 0 |
| 0 | x_9 | x_{10} | x_{11} | x_{12} | 0 |
| 0 | x_{13} | x_{14} | x_{15} | x_{16} | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 |

*

| | | |
|-------|-------|-------|
| w_1 | w_2 | w_3 |
| w_4 | w_5 | w_6 |
| w_7 | w_8 | w_9 |

=

| | | | |
|----------|----------|----------|----------|
| z_1 | z_2 | z_3 | z_4 |
| z_5 | z_6 | z_7 | z_8 |
| z_9 | z_{10} | z_{11} | z_{12} |
| z_{13} | z_{14} | z_{15} | z_{16} |

Convolution calculation on borders

| | | | | | |
|---|-----|-----|-----|-----|-----|
| 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 105 | 102 | 100 | 97 | 96 |
| 0 | 103 | 99 | 103 | 101 | 102 |
| 0 | 101 | 98 | 104 | 102 | 100 |
| 0 | 99 | 101 | 106 | 104 | 99 |
| 0 | 104 | 104 | 104 | 100 | 98 |

Image Matrix

| Kernel Matrix | | |
|---------------|----|----|
| 0 | -1 | 0 |
| -1 | 5 | -1 |
| 0 | -1 | 0 |

| | | | | |
|-----|----|-----|--|--|
| | | | | |
| 210 | 89 | 111 | | |
| | | | | |
| | | | | |
| | | | | |
| | | | | |

Output Matrix

$$\begin{aligned}
 &0 * 0 + 105 * -1 + 102 * 0 \\
 &+ 0 * -1 + 103 * 5 + 99 * -1 \\
 &+ 0 * 0 + 101 * -1 + 98 * 0 = 210
 \end{aligned}$$

<image Convolution>

Machinelearningguru.com/computer_vision/basics/convolution/image_convolution_1.html

Convolution calculation on borders

| | | | | | |
|---|-----|-----|-----|-----|-----|
| 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 105 | 102 | 100 | 97 | 96 |
| 0 | 103 | 99 | 103 | 101 | 102 |
| 0 | 101 | 98 | 104 | 102 | 100 |
| 0 | 99 | 101 | 106 | 104 | 99 |
| 0 | 104 | 104 | 104 | 100 | 98 |

Image Matrix

| Kernel Matrix | | |
|---------------|----|----|
| 0 | -1 | 0 |
| -1 | 5 | -1 |
| 0 | -1 | 0 |

| | | | | |
|-----|----|-----|--|--|
| 320 | | | | |
| 210 | 89 | 111 | | |
| | | | | |
| | | | | |
| | | | | |
| | | | | |

Output Matrix

$$\begin{aligned}
 &0 * 0 + 0 * -1 + 0 * 0 \\
 &+ 0 * -1 + 105 * 5 + 102 * -1 \\
 &+ 0 * 0 + 103 * -1 + 99 * 0 = 320
 \end{aligned}$$

<image Convolution>

Machinelearningguru.com/computer_vision/basics/convolution/image_convolution_1.html

Padding: 1 zero padding

| | | | | | |
|---|----------|----------|----------|----------|---|
| 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | X_1 | X_2 | X_3 | X_4 | 0 |
| 0 | X_5 | X_6 | X_7 | X_8 | 0 |
| 0 | X_9 | X_{10} | X_{11} | X_{12} | 0 |
| 0 | X_{13} | X_{14} | X_{15} | X_{16} | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 |

*

| | | |
|-------|-------|-------|
| W_1 | W_2 | W_3 |
| W_4 | W_5 | W_6 |
| W_7 | W_8 | W_9 |

=

| | | | |
|----------|----------|----------|----------|
| Z_1 | Z_2 | Z_3 | Z_4 |
| Z_5 | Z_6 | Z_7 | Z_8 |
| Z_9 | Z_{10} | Z_{11} | Z_{12} |
| Z_{13} | Z_{14} | Z_{15} | Z_{16} |

$$\begin{aligned}
 Z_1 &= 0 \times W_1 + 0 \times W_2 + 0 \times W_3 \\
 &\quad + 0 \times W_4 + X_1 \times W_5 + X_2 \times W_6 \\
 &\quad + 0 \times W_7 + X_5 \times W_8 + X_6 \times W_9
 \end{aligned}$$

$$\begin{aligned}
 Z_4 &= 0 \times W_1 + 0 \times W_2 + 0 \times W_3 \\
 &\quad + X_3 \times W_4 + X_4 \times W_5 + 0 \times W_6 \\
 &\quad + X_7 \times W_7 + X_8 \times W_8 + 0 \times W_9
 \end{aligned}$$

Padding: 2 zero padding

| | | | | | | | |
|---|---|----------|----------|----------|----------|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | X_1 | X_2 | X_3 | X_4 | 0 | 0 |
| 0 | 0 | X_5 | X_6 | X_7 | X_8 | 0 | 0 |
| 0 | 0 | X_9 | X_{10} | X_{11} | X_{12} | 0 | 0 |
| 0 | 0 | X_{13} | X_{14} | X_{15} | X_{16} | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |

*

| | | |
|-------|-------|-------|
| W_1 | W_2 | W_3 |
| W_4 | W_5 | W_6 |
| W_7 | W_8 | W_9 |

Stride

Input matrix

| | | | | |
|-----------------|-----------------|-----------------|-----------------|-----------------|
| X ₁ | X ₂ | X ₃ | X ₄ | X ₅ |
| X ₆ | X ₇ | X ₈ | X ₉ | X ₁₀ |
| X ₁₁ | X ₁₂ | X ₁₃ | X ₁₄ | X ₁₅ |
| X ₁₆ | X ₁₇ | X ₁₈ | X ₁₉ | X ₂₀ |
| X ₂₁ | X ₂₂ | X ₂₃ | X ₂₄ | X ₂₅ |

filter

| | |
|----------------|----------------|
| W ₁ | W ₂ |
| W ₃ | W ₄ |

*


Stride 2

Output matrix

| | |
|----------------|----------------|
| Z ₁ | Z ₂ |
| Z ₃ | Z ₄ |

Size: 2 x 2

Convolution calculation with stride 1



<understanding convolutional layers in convolutional neural networks>

[Machinelearningguru.com/computer_vision/basics/convolution/convolution_layer.html](https://machinelearningguru.com/computer_vision/basics/convolution/convolution_layer.html)

Convolution calculation with stride 2

| | | | | | | |
|---|-----|-----|-----|-----|-----|-----|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 105 | 102 | 100 | 97 | 96 | 100 |
| 0 | 103 | 99 | 103 | 101 | 102 | 100 |
| 0 | 101 | 98 | 104 | 102 | 100 | 100 |
| 0 | 99 | 101 | 106 | 104 | 99 | 100 |
| 0 | 104 | 104 | 104 | 100 | 98 | 100 |

Image Matrix

| | | |
|---------------|----|----|
| Kernel Matrix | | |
| 0 | -1 | 0 |
| -1 | 5 | -1 |
| 0 | -1 | 0 |

| | | | | |
|-----|-----|-----|--|--|
| 320 | 198 | 182 | | |
| | | | | |
| | | | | |
| | | | | |
| | | | | |
| | | | | |

Output Matrix

$$\begin{aligned}
 &0 * 0 + 0 * -1 + 0 * 0 \\
 &+ 97 * -1 + 96 * 5 + 99 * -1 \\
 &+ 101 * 0 + 102 * -1 + 101 * 0 = 182
 \end{aligned}$$

<understanding convolutional layers in convolutional neural networks>

[Machinelearningguru.com/computer_vision/basics/convolution/convolution_layer.html](https://machinelearningguru.com/computer_vision/basics/convolution/convolution_layer.html)


Summary of convolutions

- $n \times n$ image
- $f \times f$ filter
- padding p
- stride s

Output size:

$$\left\lceil \frac{n + 2p - f}{s} + 1 \right\rceil \times \left\lceil \frac{n + 2p - f}{s} + 1 \right\rceil$$


Convolutions on RGB images



<understanding convolutional layers in convolutional neural networks>

[Machinelearningguru.com/computer_vision/basics/convolution/convolution_layer.html](https://machinelearningguru.com/computer_vision/basics/convolution/convolution_layer.html)


Multiple filters (channels)



<deep learning, Andrew Ng>

Multiple filters (channels)

$$Z = W * X + b$$
$$a = \text{ReLU}(Z)$$



<deep learning, Andrew Ng>

Number of parameters in one layer

If there are 10 filters that have the size of $3 \times 3 \times 3$ in one layer of a neural network, how many parameters does that layer have?

Number of parameters in one layer

If there are 10 filters that are $3 \times 3 \times 3$ in one layer of a neural network, how many parameters does that layer have?

For 1 filter:

$3 \times 3 \times 3 = 27$ for W

1 for bias

For 10 filters:

Then, $28 * 10 = 280$

Summary of size in matrix

If layer l is a convolution layer:

$f^{[l]}$ = filter size

$p^{[l]}$ = padding

$s^{[l]}$ = stride

Input size:

$$n_H^{[l-1]} \times n_W^{[l-1]} \times n_C^{[l-1]}$$

output size:

$$n_H^{[l]} \times n_W^{[l]} \times n_C^{[l]}$$

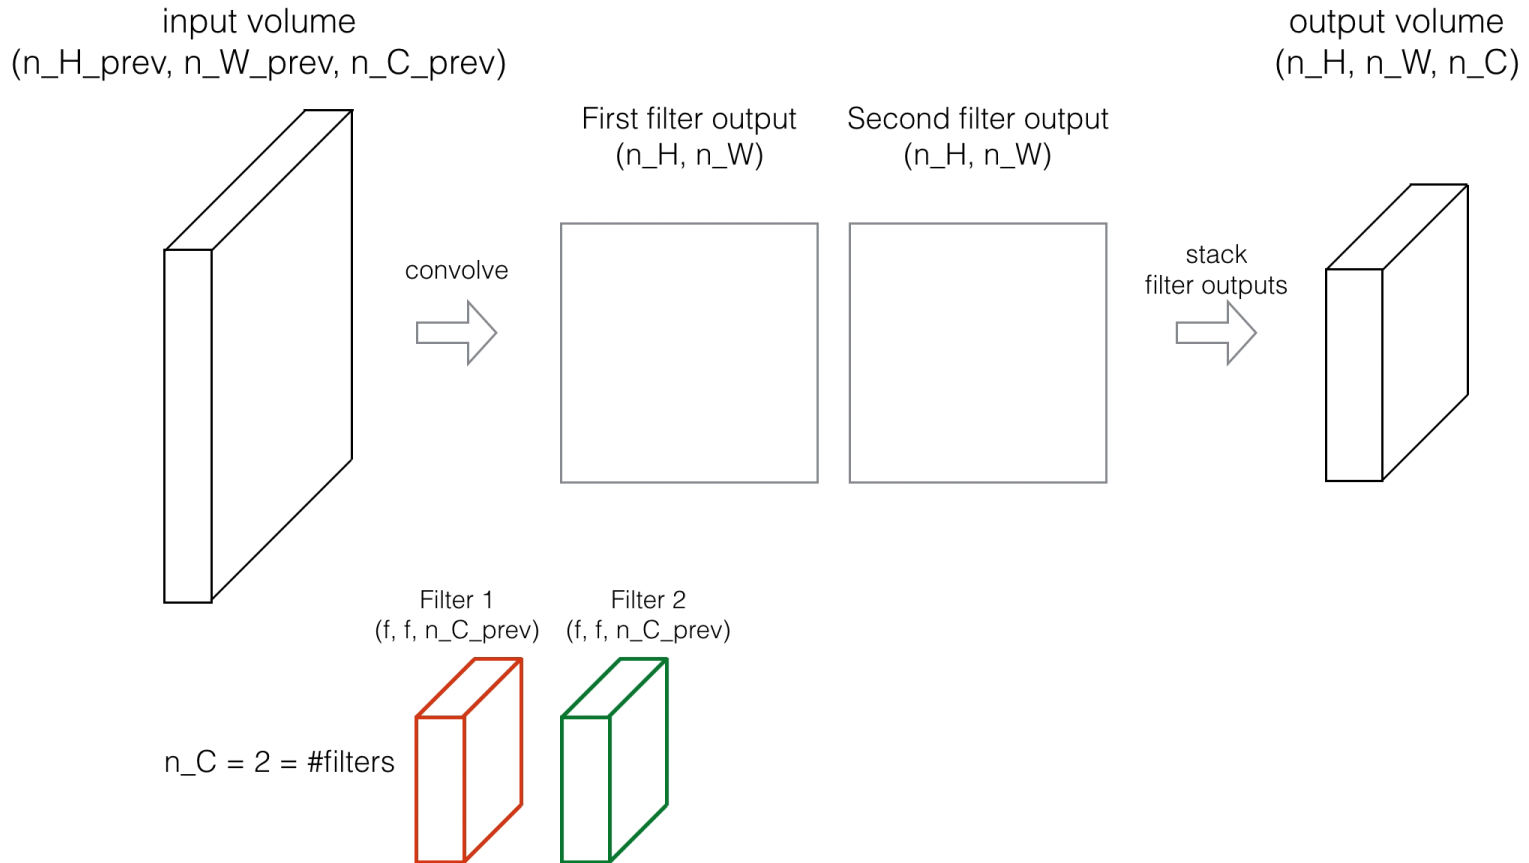
$$n_H^{[l]} = \frac{n_H^{[l-1]} + 2p^{[l]} - f^{[l]}}{s^{[l]}} + 1$$

$$n_W^{[l]} = \frac{n_W^{[l-1]} + 2p^{[l]} - f^{[l]}}{s^{[l]}} + 1$$

Weights: $f^{[l]} \times f^{[l]} \times n_C^{[l-1]} \times n_C^{[l]}$

bias: $1 \times 1 \times 1 \times n_C^{[l]}$

How do convolutions work?



Summary

- Compare CNN vs. DNN
- 2D image filter
- Convolution calculation
- Convolution on RGB images
- Multiple filters
- Size of matrix in convolution layers