



CMPE 258, Deep Learning

Sequence learning & NLP

April 26, 2018

DMH 149A

Taehee Jeong

Ph.D., Data Scientist

Group Project schedule

Presentation date : 5/8, 5/10

Report (including code) due date : 5/6

Number of members : 1 to 4

Content: DNN, CNN, RNN related

Platform : Pandas, Numpy, tensorflow, keras (please discuss with me for others)

Grading policy:

- Content : 40 pts

- ; Creativity in data collection, Neural network architecture / algorithm, application (same quality as a conference paper)

- Presentation : 20 pts

- Report : 20 pts


- Code : 20 pts

Applications of Image analysis using Deep Learning

- Medical Image Analysis
: MRI, CT, Ultrasound, X-ray
- Self-driving car / drone
- Home security
- Care for the Disable
- Cloth searching

Sequential data

Text, Video, and Audio



Time series: finance, industry, medicine...



Coursera: Introduction to Deep Learning, National Research University Higher School of Economics

Applications of Recurrent Neural Networks

Speech recognition



"The quick brown fox jumped over the lazy dog."

Music generation

∅



Sentiment classification

"There is nothing to like in this movie."



DNA sequence analysis

AGCCCCTGTGAGGAACTAG



AGCCCCTGTGAGGAACTAG


Machine translation

Voulez-vous chanter avec moi?



Do you want to sing with me?

Video activity recognition



Running

Name entity recognition

Yesterday, Harry Potter met Hermione Granger.



Yesterday, **Harry Potter** met **Hermione Granger**.

Coursera: Deep learning Specialization, Andrew Ng

Speech recognition



Amazon Echo
(Alexa)



Baidu DuerOS
(xiaodunihao)



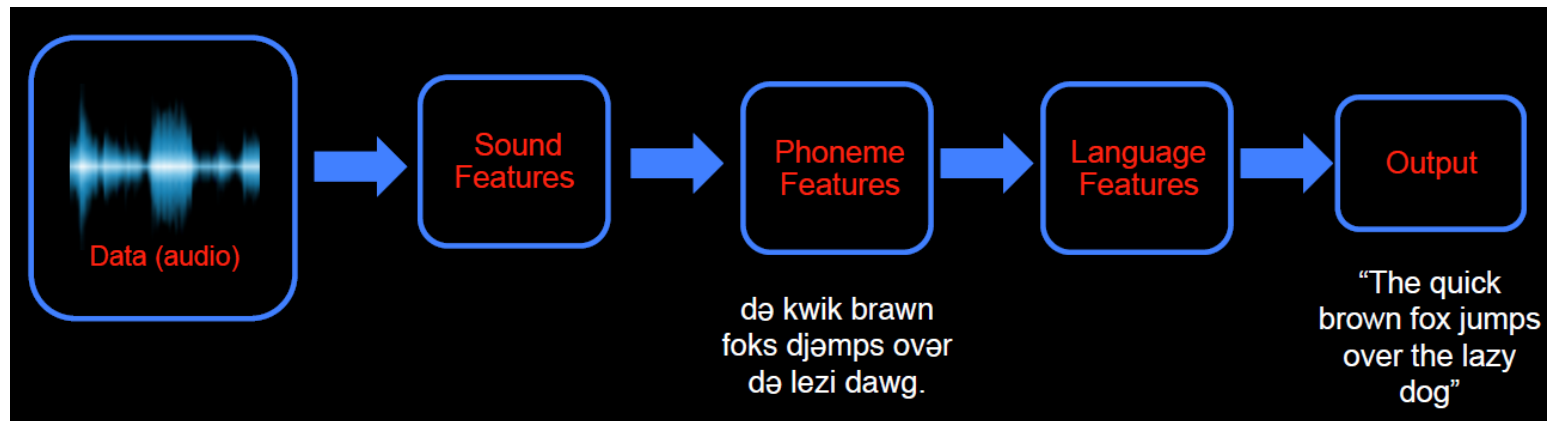
Apple Siri
(Hey Siri)



Google Home
(Okay Google)

Coursera: Deep learning Specialization, Andrew Ng


Speech recognition



GTC 2015, Andrew Ng

Deep speech


Baidu research



GTC 2015, Andrew Ng

Applications of Speech recognition

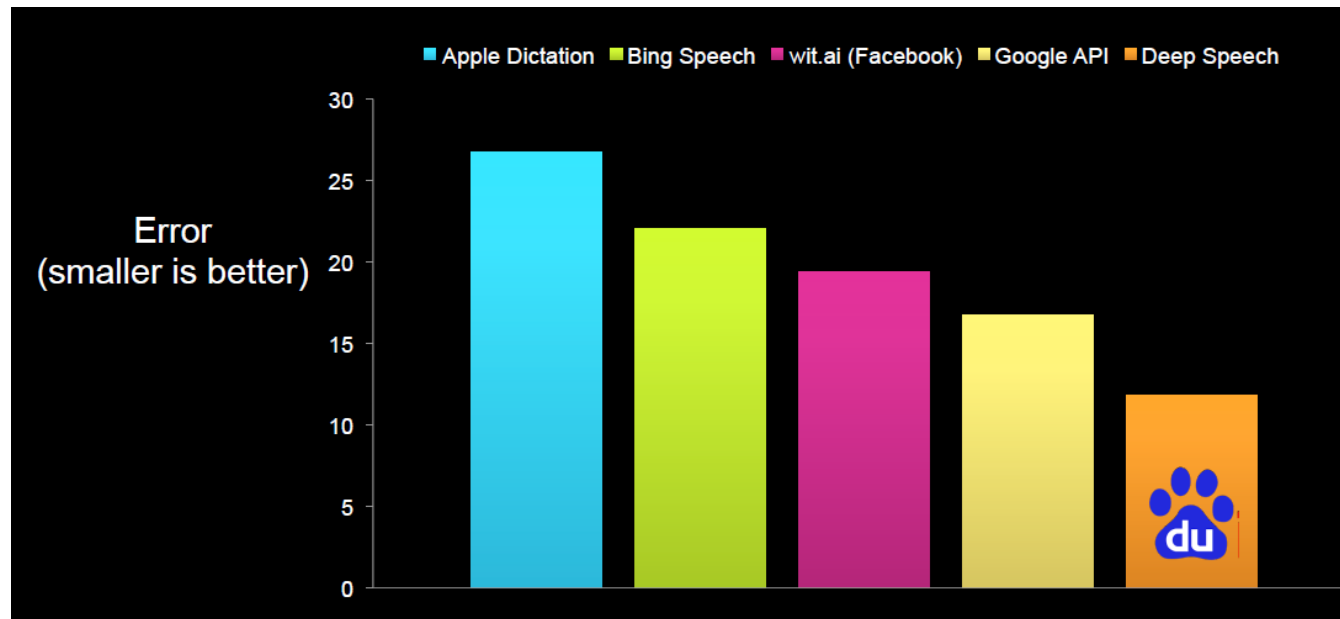
- Car interfaces
- Home appliances
- Wearables
- Care for the Disable



GTC 2015, Andrew Ng

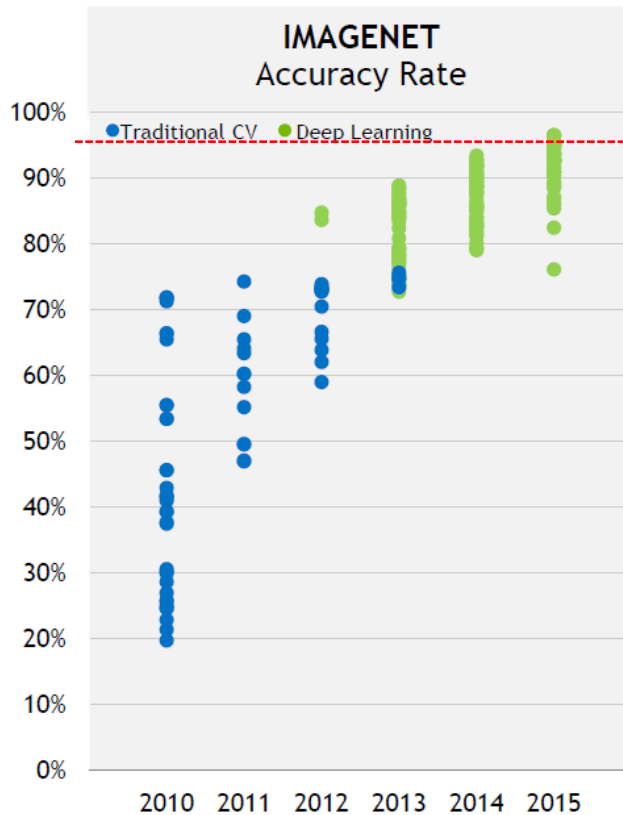
Speech recognition errors

There are room to improve compared to image recognition.



GTC 2015, Andrew Ng

Image classification

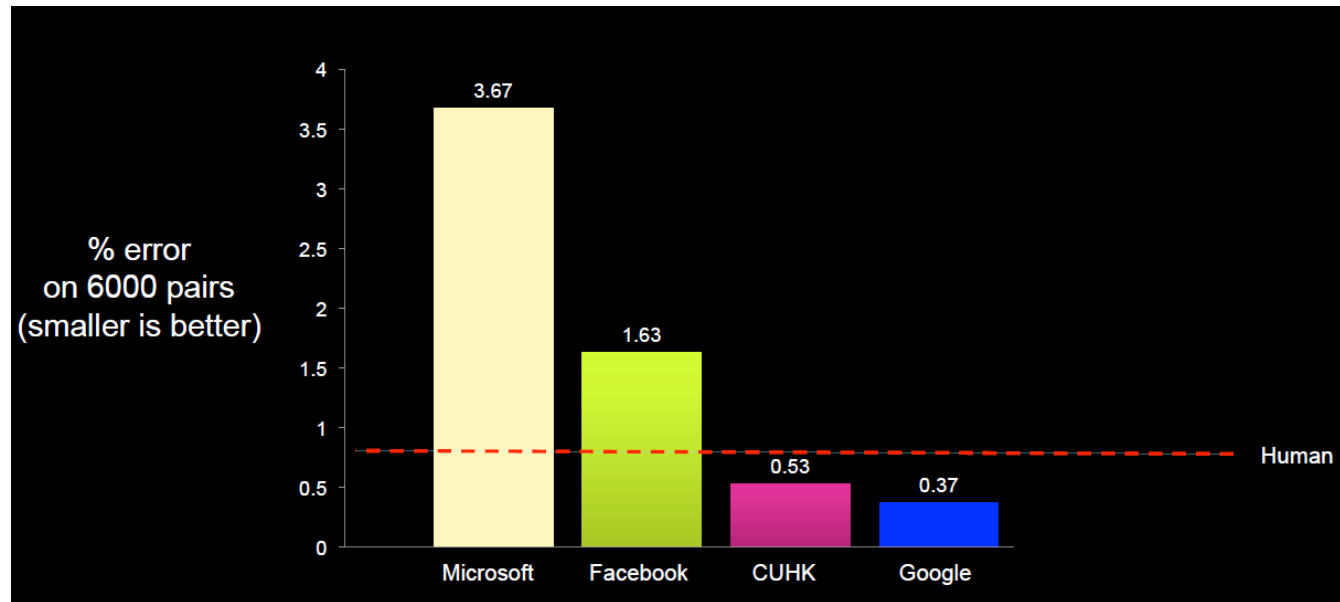


Top 5 error rate for image classification in ILSVRC
ImageNet challenge
From 26% to 3% in six years

Human accuracy : 95 ~ 97 %

Image-Net Large-Scale Visual Recognition Challenge (ILSVRC)
<http://image-net.org/challenges/LSVRC/>
<https://www.slideshare.net/NVIDIA/nvidia-ces-2016-press-conference>

Face recognition errors



GTC 2015, Andrew Ng

Music generation by RNN

Chorus 1
Intro-Verse
Style: swing
Giant Steps
John Coltrane

13

Chorus 1
Intro-Verse
Style: swing
Now's The Time
Charlie Parker

12

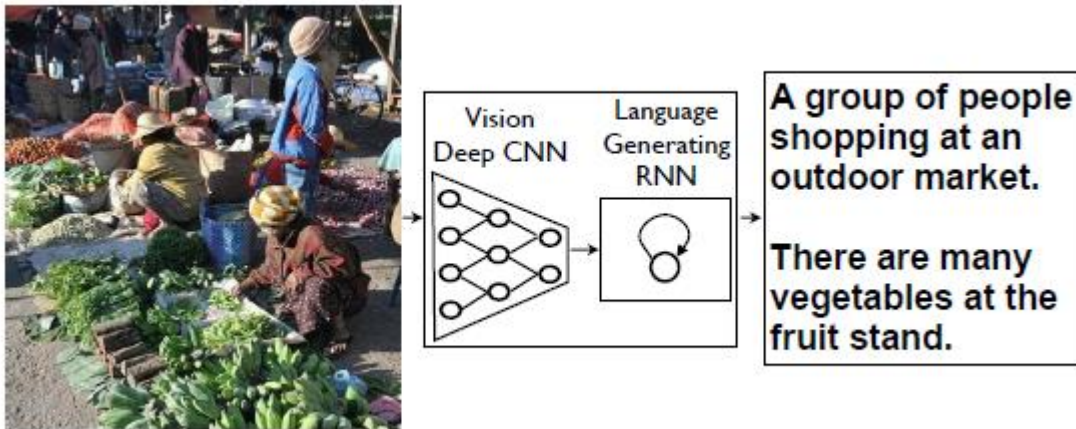
LEARNING JAZZ GRAMMARS, SMC09, Jon Gillick et al.

Text generation using RNN

They had no choice but the most recent univerebeen fairly uncomfortable and dangerous as ever. As long as he dived experience that it was not uncertain that even Harry had taken in black tail as the train roared and was thin, but Harry, Ron, and Hermione, at the fact that he was in complete disarraying the rest of the class holding him, he should have been able to prove them.

<https://chunml.github.io/ChunML.github.io/project/Creating-Text-Generator-Using-Recurrent-Neural-Network/>

A Neural Image Caption Generator



Show and Tell: A Neural Image Caption Generator, Oriol Vinyals et al., 2015,

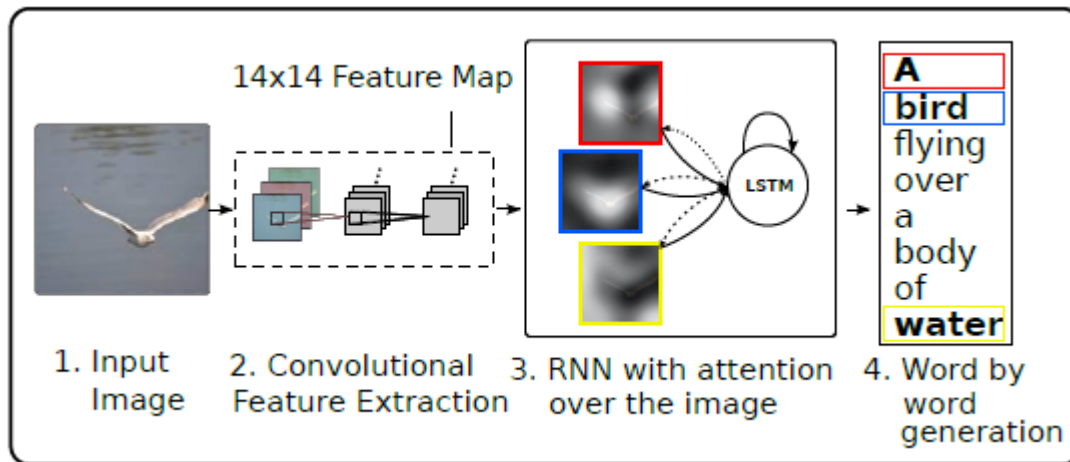
A Neural Image Caption Generator



Show and Tell: A Neural Image Caption Generator, Oriol Vinyals et al., 2015,

Neural Image Caption Generation with Visual Attention

Visualized attentional maps



Show, Attend and Tell: Neural Image Caption Generation with Visual Attention, Kelvin Xu et al., 2016

Neural Image Caption Generation with Visual Attention



A woman is throwing a frisbee in a park.



A dog is standing on a hardwood floor.



A stop sign is on a road with a mountain in the background.



A little girl sitting on a bed with a teddy bear.



A group of people sitting on a boat in the water.



A giraffe standing in a forest with trees in the background.

Show, Attend and Tell: Neural Image Caption Generation with Visual Attention, Kelvin Xu et al., 2016

Natural Language Process (NLP)

- Speech recognition
- Machine translation
- Chatbots (question answering)
- Sentiment classification
- Name entity recognition

Approach for NLP

- Machine Learning
- Convolutional Neural Network
- Recurrent Neural Network

Sentiment analysis

Text is a sequence of words

Word is a sequence of characters

How to separate words from a sentence?

Tokenization

Tokenization is a process that splits an input sequence into tokens.

We can split token by space, punctuation, a set of rule.

Tokenization

In Düsseldorf I took my hat off. But I can't put it back on.

WhitespaceTokenizer

1.

In Düsseldorf I took my hat off.

2.

But I can't put it back on.

WordPunctTokenizer

1.

In Düsseldorf I took my hat off .

2.

But I can ' t put it back on .

TreebankWordTokenizer

1.

In Düsseldorf I took my hat off .

2.

But I ca n't put it back on .

<http://text-processing.com/demo/tokenize/>

Python tokenization example

```
import nltk  
text = "This is Andrew's text, isn't it?"
```

```
tokenizer = nltk.tokenize.WhitespaceTokenizer()  
tokenizer.tokenize(text)
```

```
['This', 'is', "Andrew's", 'text,', "isn't", 'it?']
```

```
tokenizer = nltk.tokenize.TreebankWordTokenizer()  
tokenizer.tokenize(text)
```

```
['This', 'is', 'Andrew', "'s", 'text', ',', 'is', "n't",  
'it', '?']
```

```
tokenizer = nltk.tokenize.WordPunctTokenizer()  
tokenizer.tokenize(text)
```

```
['This', 'is', 'Andrew', "'", 's', 'text', ',', 'isn',  
"'", 't', 'it', '?']
```

Coursera: Natural Language Processing, National Research University Higher School of Economics

Token normalization

Same token for different forms of words

- Examples
 - wolf, wolves → wolf
 - talk, talks → talk
- Stemming
 - removes and replaces suffixes to get to the root form of a word, which is called as stem.
- Lemmatization
 - returns the base or dictionary form of a word, which is known as lemma.

Coursera: Natural Language Processing, National Research University Higher School of Economics

Python stemming example

```
import nltk
text = "feet cats wolves talked"
tokenizer = nltk.tokenize.TreebankWordTokenizer()
tokens = tokenizer.tokenize(text)
```

```
stemmer = nltk.stem.PorterStemmer()
" ".join(stemmer.stem(token) for token in tokens)
```

```
u'feet cat wolv talk'
```

```
stemmer = nltk.stem.WordNetLemmatizer()
" ".join(stemmer.lemmatize(token) for token in tokens)
```


```
u'foot cat wolf talked'
```

Transforming tokens into features

Bag of words (BOW)

For each token, we have a feature column, which is called text vectorization.

good movie	
not a good movie	
did not like	




good	movie	not	a	did	like
1	1	0	0	0	0
1	1	1	1	0	0
0	0	1	0	1	1

Coursera: Natural Language Processing, National Research University Higher School of Economics

Preserve some ordering

N-grams: Token pairs, triplets, etc.

good movie		good movie	movie	did not	a	...
not a good movie		1	1	0	0	...
did not like		1	1	0	1	...
		0	0	1	0	...

Coursera: Natural Language Processing, National Research University Higher School of Economics

Remove some n-grams

- High frequency n-grams
 - Articles, prepositions, etc. (example: and, a, the)
 - They are called stop-words. They do not help to discriminate texts.
- Low frequency n-grams
 - Typos, rare words

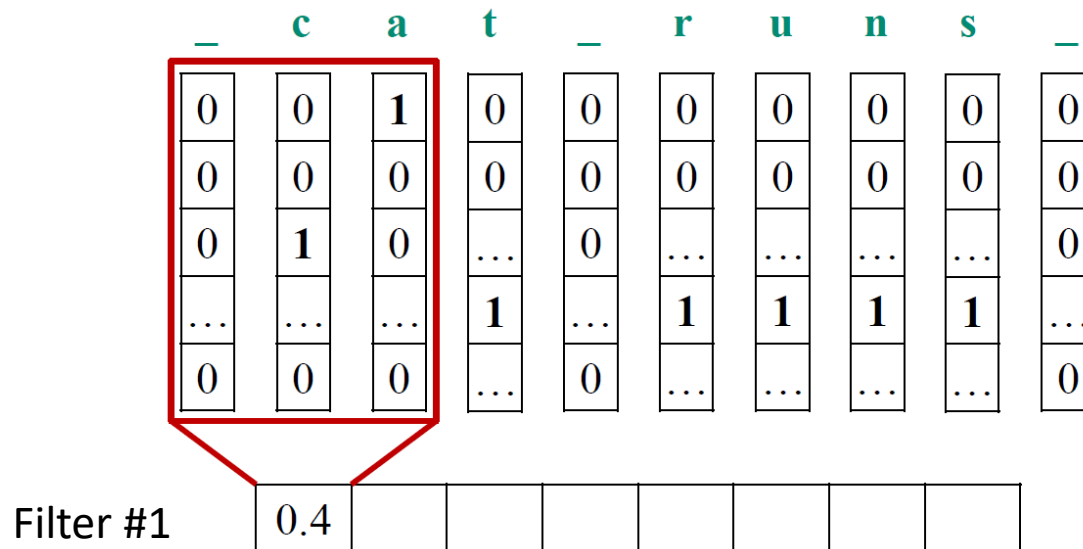
Text as a sequence of characters

One-hot encoding characters, length ~ 70

_	c	a	t	_	r	u	n	s	_
0	0	1	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	1	0	...	0	0
...	1	...	1	1	1	1	...
0	0	0	...	0	0

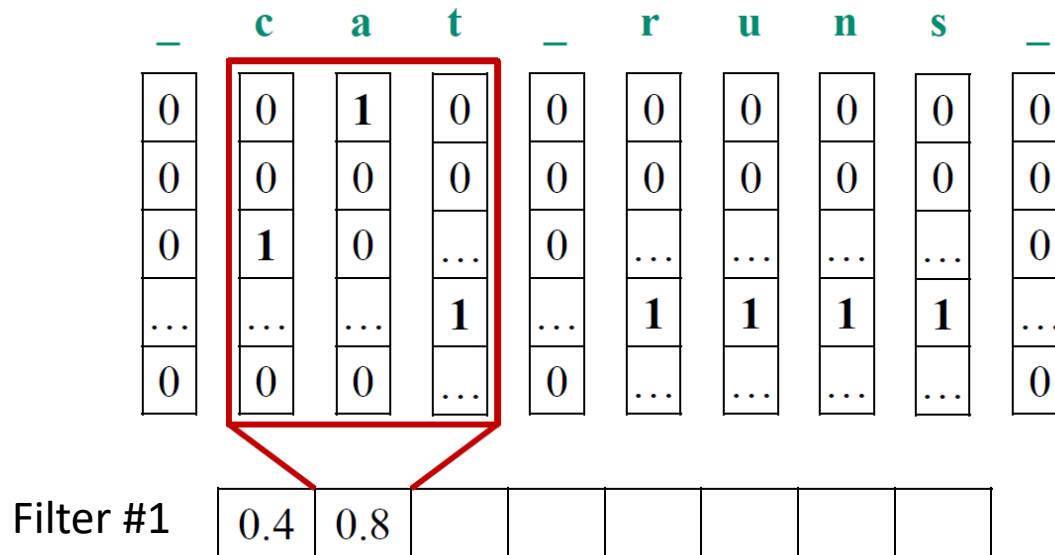
Coursera: Natural Language Processing, National Research University Higher School of Economics

1D convolutions on characters

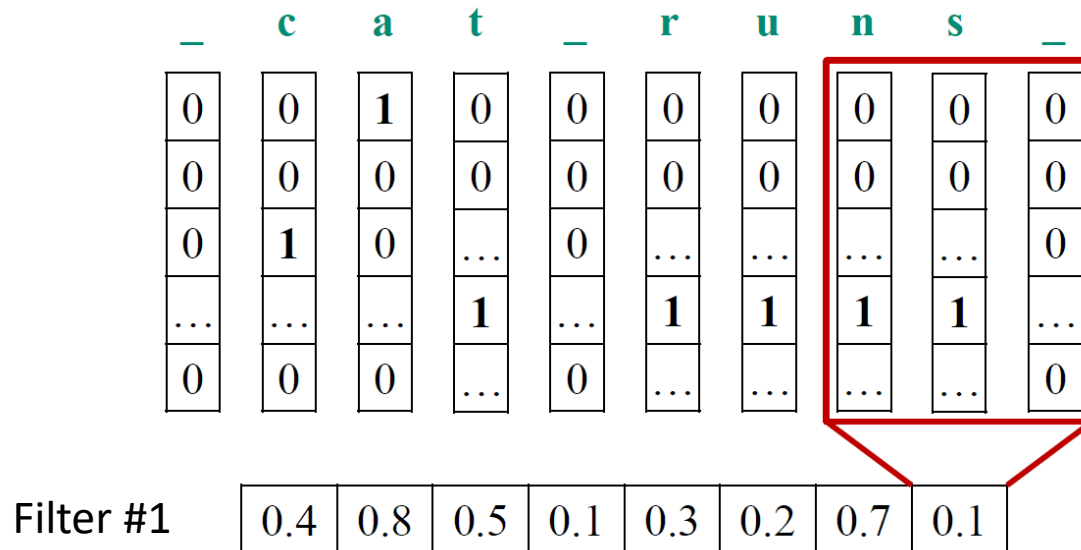


Coursera: Natural Language Processing, National Research University Higher School of Economics

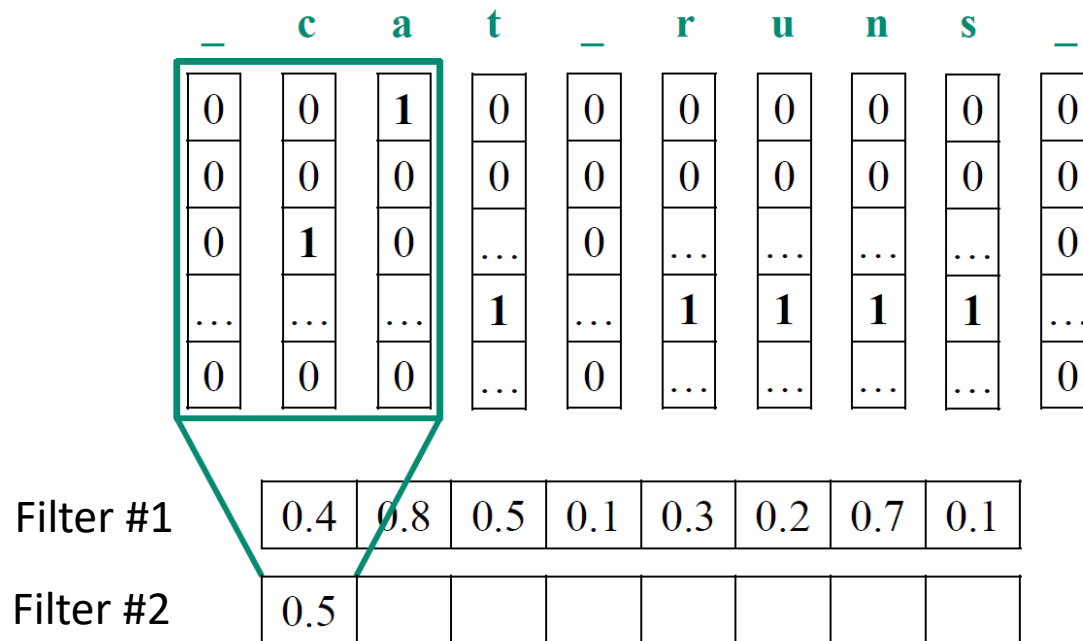
1D convolutions on characters



1D convolutions on characters

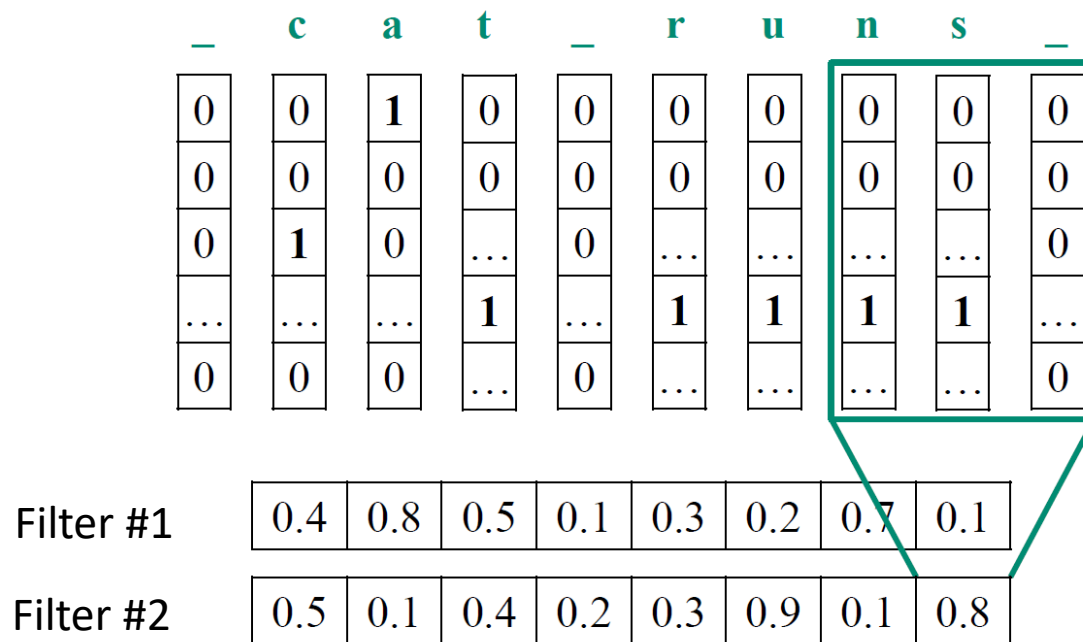


1D convolutions on characters



Coursera: Natural Language Processing, National Research University Higher School of Economics

1D convolutions on characters



Coursera: Natural Language Processing, National Research University Higher School of Economics

1D convolutions on characters

_	c	a	t	_	r	u	n	s	_
0	0	1	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	1	0	...	0	0
...	1	...	1	1	1	1	...
0	0	0	...	0	0

Filter #1	0.4	0.8	0.5	0.1	0.3	0.2	0.7	0.1
-----------	-----	-----	-----	-----	-----	-----	-----	-----


Filter #2	0.5	0.1	0.4	0.2	0.3	0.9	0.1	0.8
-----------	-----	-----	-----	-----	-----	-----	-----	-----

Filter #3	0.4	0.7	0.3	0.7	0.5	0.5	0.9	0.4
-----------	-----	-----	-----	-----	-----	-----	-----	-----

Coursera: Natural Language Processing, National Research University Higher School of Economics

Max pooling

Provides position invariance for character n-gram



Coursera: Natural Language Processing, National Research University Higher School of Economics

Max pooling

Provides position invariance for character n-gram

Filter #1

0.4	0.8	0.5	0.1	0.3	0.2	0.7	0.1
-----	-----	-----	-----	-----	-----	-----	-----

Filter #2

0.5	0.1	0.4	0.2	0.3	0.9	0.1	0.8
-----	-----	-----	-----	-----	-----	-----	-----

Filter #3

0.4	0.7	0.3	0.7	0.5	0.5	0.9	0.4
-----	-----	-----	-----	-----	-----	-----	-----

0.8	0.5		
-----	-----	--	--

--	--	--	--

--	--	--	--

Coursera: Natural Language Processing, National Research University Higher School of Economics

Max pooling

Provides position invariance for character n-gram

Filter #1

0.4	0.8	0.5	0.1	0.3	0.2	0.7	0.1
-----	-----	-----	-----	-----	-----	-----	-----

Filter #2

0.5	0.1	0.4	0.2	0.3	0.9	0.1	0.8
-----	-----	-----	-----	-----	-----	-----	-----

Filter #3

0.4	0.7	0.3	0.7	0.5	0.5	0.9	0.4
-----	-----	-----	-----	-----	-----	-----	-----

0.8	0.5	0.3	0.7
-----	-----	-----	-----

--	--	--	--

--	--	--	--

Coursera: Natural Language Processing, National Research University Higher School of Economics

Max pooling

Provides position invariance for character n-gram

Filter #1

0.4	0.8	0.5	0.1	0.3	0.2	0.7	0.1
-----	-----	-----	-----	-----	-----	-----	-----

Filter #2

0.5	0.1	0.4	0.2	0.3	0.9	0.1	0.8
-----	-----	-----	-----	-----	-----	-----	-----

Filter #3

0.4	0.7	0.3	0.7	0.5	0.5	0.9	0.4
-----	-----	-----	-----	-----	-----	-----	-----


Pooling
output

0.8	0.5	0.3	0.7
-----	-----	-----	-----

0.5	0.4	0.9	0.8
-----	-----	-----	-----

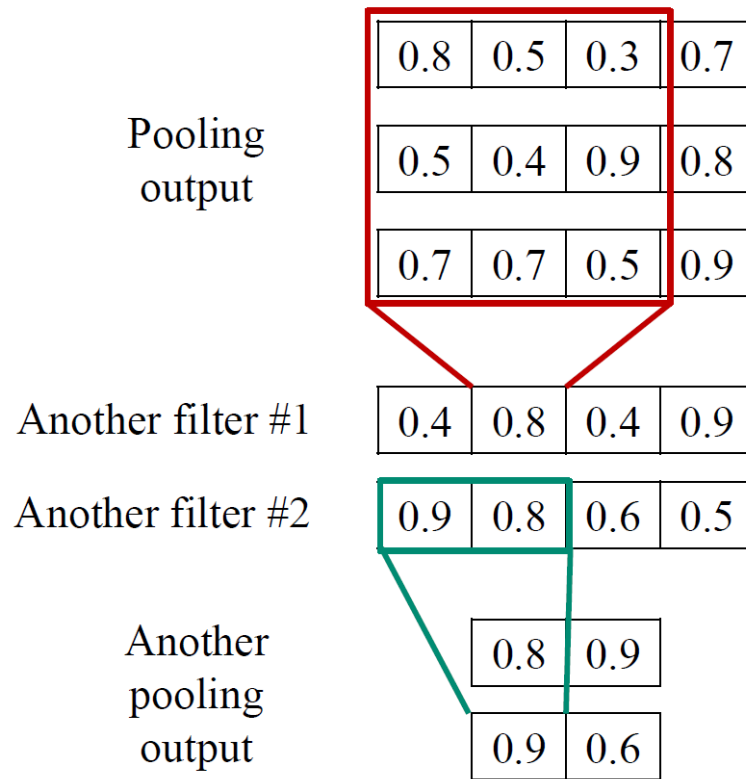
0.7	0.7	0.5	0.9
-----	-----	-----	-----

Repeat 1D convolutions + pooling




Coursera: Natural Language Processing, National Research University Higher School of Economics

Repeat 1D convolutions + pooling



Coursera: Natural Language Processing, National Research University Higher School of Economics

Repeat 1D convolutions + pooling



Coursera: Natural Language Processing, National Research University Higher School of Economics

Example of 1D convolution architecture for characters

- Characters of text : 1014
- Apply 1D convolution + max pooling 6 times
- Kernels widths : 7,7,3,3,3,3
- # of Filters at each step : 1024
- Output: 1024 x 34 matrix of features
- Apply multiclass classification for sentiment analysis

Coursera: Natural Language Processing, National Research University Higher School of Economics

Summary

- Sequence data and its application using RNN
- Natural Language Process
- Tokenization
- Bag of Words, Word vector
- 1D convolution with characters