



CMPE 258, Deep Learning

Sequence learning & NLP

May 01, 2018

DMH 149A

Taehee Jeong

Ph.D., Data Scientist

Group Project schedule

Presentation date : 5/8, 5/10

Report (including code) due date : 5/6

Number of members : 1 to 4

Content: DNN, CNN, RNN related

Platform : Pandas, Numpy, tensorflow, keras (please discuss with me for others)

Grading policy:

- Content : 40 pts


- ; Creativity in data collection, Neural network architecture / algorithm, application (same quality as a conference paper)

- Presentation : 20 pts

- Report : 20 pts

- Code : 20 pts

Mid-term exam 2 score




Median = 93.33

Mean = 87.67

The submitted ipynb should be executable without any extra work and supposed be finished within 60 minutes. If extra effort is needed to get reasonable result (whatever it is), 5 to 20 points for each event will be deducted.

Score = accuracy of testing data -20 + 10 (print out of CNN architecture) + 10 (plot of cost versus number of iteration)

| Layer | Type | Size | Channels | Kernel size | Stride | Padding | Function |
|-------|----------------------|---------|----------|-------------|--------|---------|----------|
| 0 | Input | 64 x 64 | 3 | | | | |
| 1 | Convolution (C1) | 32 x 32 | 8 | 4 x 4 | 2 | 1 | ReLU |
| 1 | Pooling (P1) | 28 x 28 | 8 | 5 x 5 | 1 | 0 | max |
| 2 | Convolution (C2) | 13 x 13 | 16 | 4 x 4 | 2 | 0 | ReLU |
| 2 | Pooling (P2) | 9 x 9 | 16 | 5 x 5 | 1 | 0 | Avg |
| 3 | Flatten (F3) | 1296 | | | | | |
| 4 | Fully connected (F4) | 108 | | | | | ReLU |
| 5 | Fully connected (F5) | 6 | | | | | Sigmoid |



Applications of Recurrent Neural Networks

Speech recognition



"The quick brown fox jumped over the lazy dog."

Music generation

∅



Sentiment classification

"There is nothing to like in this movie."



DNA sequence analysis

AGCCCCTGTGAGGAACTAG



AGCCCCTGTGAGGAACTAG


Machine translation

Voulez-vous chanter avec moi?



Do you want to sing with me?

Video activity recognition



Running

Name entity recognition

Yesterday, Harry Potter met Hermione Granger.



Yesterday, **Harry Potter** met **Hermione Granger**.

Coursera: Deep learning Specialization, Andrew Ng

Natural Language Process (NLP)

- Speech recognition
- Machine translation
- Chatbots (question answering)
- Sentiment classification
- Name entity recognition

Sentiment analysis

- Text is a sequence of words
- Word is a sequence of characters

How to separate words from a sentence?

Tokenization

- Tokenization is a process that splits an input sequence into tokens.
- We can split token by space, punctuation, a set of rule.

Python tokenization example

```
import nltk  
text = "This is Andrew's text, isn't it?"
```

```
tokenizer = nltk.tokenize.WhitespaceTokenizer()  
tokenizer.tokenize(text)
```

```
['This', 'is', "Andrew's", 'text,', "isn't", 'it?']
```

```
tokenizer = nltk.tokenize.TreebankWordTokenizer()  
tokenizer.tokenize(text)
```

```
['This', 'is', 'Andrew', "'s", 'text', ',', 'is', "n't",  
'it', '?']
```

```
tokenizer = nltk.tokenize.WordPunctTokenizer()  
tokenizer.tokenize(text)
```

```
['This', 'is', 'Andrew', "'", 's', 'text', ',', 'isn',  
"'", 't', 'it', '?']
```

Coursera: Natural Language Processing, National Research University Higher School of Economics

Token normalization

Same token for different forms of words

- Examples
 - wolf, wolves → wolf
 - talk, talks → talk
- Stemming
 - removes and replaces suffixes to get to the root form of a word, which is called as stem.
- Lemmatization
 - returns the base or dictionary form of a word, which is known as lemma.

Coursera: Natural Language Processing, National Research University Higher School of Economics

Python stemming example

```
import nltk
text = "feet cats wolves talked"
tokenizer = nltk.tokenize.TreebankWordTokenizer()
tokens = tokenizer.tokenize(text)
```

```
stemmer = nltk.stem.PorterStemmer()
" ".join(stemmer.stem(token) for token in tokens)
```

```
u'feet cat wolv talk'
```

```
stemmer = nltk.stem.WordNetLemmatizer()
" ".join(stemmer.lemmatize(token) for token in tokens)
```


```
u'foot cat wolf talked'
```

Transforming tokens into features

Bag of words (BOW)

For each token, we have a feature column, which is called text vectorization.

| | |
|------------------|--|
| good movie | |
| not a good movie | |
| did not like | |




| good | movie | not | a | did | like |
|------|-------|-----|---|-----|------|
| 1 | 1 | 0 | 0 | 0 | 0 |
| 1 | 1 | 1 | 1 | 0 | 0 |
| 0 | 0 | 1 | 0 | 1 | 1 |

Coursera: Natural Language Processing, National Research University Higher School of Economics

Preserve some ordering

N-grams: Token pairs, triplets, etc.

| | | | | | | |
|------------------|---|---------------|-------|---------|---|-----|
| good movie |  | good movie | movie | did not | a | ... |
| not a good movie | | 1 | 1 | 0 | 0 | ... |
| did not like | | 1 | 1 | 0 | 1 | ... |
| | | 0 | 0 | 1 | 0 | ... |

Coursera: Natural Language Processing, National Research University Higher School of Economics

Remove some n-grams

- High frequency n-grams
 - Articles, prepositions, etc. (example: and, a, the)
 - They are called stop-words. They do not help to discriminate texts.
- Low frequency n-grams
 - Typos, rare words

Word Embedding

- Convert texts into numbers
- Map a word to a vector using a dictionary
- Applications
 - Sentiment analysis of reviews (amazon, movie review)
 - Document or news classification or clustering (google)

<https://www.analyticsvidhya.com/blog/2017/06/word-embeddings-count-word2veec/>

Frequency based Embedding

- Count Vector
 - Frequency : Number of times a word has appeared in the document
 - Presence : Has the word appeared in the document?
- TF-IDF Vector
- Co-occurrence Vector

<https://www.analyticsvidhya.com/blog/2017/06/word-embeddings-count-word2veec/>

Word Matrix

| | Document 1 | Document 2 | Document 3 | Document 4 | Document 5 | Document 6 | Document 7 | Document 8 |
|-----------|------------|------------|------------|------------|------------|------------|------------|------------|
| Term(s) 1 | 10 | 0 | 1 | 0 | 0 | 0 | 0 | 2 |
| Term(s) 2 | 0 | 2 | 0 | 0 | 0 | 18 | 0 | 2 |
| Term(s) 3 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 2 |
| Term(s) 4 | 6 | 0 | 0 | 4 | 6 | 0 | 0 | 0 |
| Term(s) 5 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 2 |
| Term(s) 6 | 0 | 0 | 1 | 0 | 0 | 1 | 0 | 0 |
| Term(s) 7 | 0 | 1 | 8 | 0 | 0 | 0 | 0 | 0 |
| Term(s) 8 | 0 | 0 | 0 | 0 | 0 | 3 | 0 | 0 |

Word Vector (Passage Vector) →

Document Vector ↗

How to make Term (word) features

1. All words in a dictionary
2. Unique words in corpus(all documents)

<https://www.analyticsvidhya.com/blog/2017/06/word-embeddings-count-word2veec/>

TF-IDF

- N-gram with smaller frequency can be more discriminating because it can capture a specific issue in the text
- Term frequency (TF)
 - Frequency for term (or n-gram) t in document d

| | |
|----------------|--------------------------------------|
| term frequency | $f_{t,d} / \sum_{t' \in d} f_{t',d}$ |
|----------------|--------------------------------------|

- Inverse document frequency (IDF)
 - $N = D$: total number of documents in corpus
 - $\{d \in D : t \in d\}$: Number of documents where the term t appears

| |
|---|
| $\text{idf}(t, D) = \log \frac{N}{ \{d \in D : t \in d\} }$ |
|---|


Coursera: Natural Language Processing, National Research University Higher School of Economics

TF-IDF

$$\text{Tfidf}(t,d,D) = \text{tf}(t,d) \times \text{idf}(t,D)$$

A high weight in TF-IDF means a high term frequency (in a given document) and a low document frequency of the term in a all collection of documents

| | |
|------------------|--|
| good movie | |
| not a good movie | |
| did not like | |



| good movie | movie | did not | ... |
|------------|-------|---------|-----|
| 0.17 | 0.17 | 0 | ... |
| 0.17 | 0.17 | 0 | ... |
| 0 | 0 | 0.47 | ... |

Coursera: Natural Language Processing, National Research University Higher School of Economics

Term frequency & Inverse document frequency

1-gram

| | | | | term frequency | | | | | | |
|------------------|------|-------|-----|----------------------------|-----|-----|------|---|----|-----|
| text | good | movie | not | a | did | not | like | I | it | one |
| good movie | | | | | | | | | | |
| not a good movie | | | | | | | | | | |
| did not like | | | | | | | | | | |
| I like it | | | | | | | | | | |
| good one | | | | | | | | | | |
| | | | | inverse document frequency | | | | | | |
| text | good | movie | not | a | did | not | like | I | it | one |
| good movie | | | | | | | | | | |
| not a good movie | | | | | | | | | | |
| did not like | | | | | | | | | | |
| I like it | | | | | | | | | | |
| good one | | | | | | | | | | |

Python TF-IDF example

```
from sklearn.feature_extraction.text import TfidfVectorizer
import pandas as pd
texts = [
    "good movie", "not a good movie", "did not like",
    "i like it", "good one"
]
tfidf = TfidfVectorizer(min_df=2, max_df=0.5, ngram_range=(1, 2))
features = tfidf.fit_transform(texts)
pd.DataFrame(
    features.todense(),
    columns=tfidf.get_feature_names()
)
```

| | good movie | like | movie | not |
|---|------------|----------|----------|----------|
| 0 | 0.707107 | 0.000000 | 0.707107 | 0.000000 |
| 1 | 0.577350 | 0.000000 | 0.577350 | 0.577350 |
| 2 | 0.000000 | 0.707107 | 0.000000 | 0.707107 |
| 3 | 0.000000 | 1.000000 | 0.000000 | 0.000000 |
| 4 | 0.000000 | 0.000000 | 0.000000 | 0.000000 |

Coursera: Natural Language Processing, National Research University Higher School of Economics

Problem of one-hot representation


- It treats each word as a thing unto itself, and it doesn't allow an algorithm to easily generalize the cross words.
- Example
 - Sally Johnson is an orange farmer.
 - Robert Lin is an apple farmer.
 - Robert Lin is a durian cultivator.
- product of any two word vector is zero.

Featured representation: word embedding

| | Man | Woman | King | Queen | Apple | Orange |
|--------|------|-------|-------|-------|-------|--------|
| Gender | -1 | 1 | -0.95 | 0.97 | 0 | 0.01 |
| Royal | 0.01 | 0.02 | 0.93 | 0.95 | -0.01 | 0 |
| Age | 0.03 | 0.02 | 0.7 | 0.69 | 0.03 | -0.02 |
| Food | 0.04 | 0.01 | 0.02 | 0.01 | 0.95 | 0.97 |
| Size | | | | | | |
| Cost | | | | | | |
| alive | | | | | | |
| verb | | | | | | |

Word embedding vector

Deep face feature vector



If $x^{(i)}, x^{(j)}$ are same person, $d(f(x^{(i)}) - f(x^{(j)}))$ is small.

If $x^{(i)}, x^{(j)}$ are different persons, $d(f(x^{(i)}) - f(x^{(j)}))$ is large.

Taigman et. al., 2014. DeepFace closing the gap to human level performance


Transfer learning and word embeddings

- Learn word embeddings from large text corpus. (1-100B words)
(Or download pre-trained embedding online.)
- Transfer embedding to new task with smaller training set. (say, 100k words)
- Continue to fine-tune the word embeddings with new data.

Coursera: Deep learning Specialization, Andrew Ng

Co-occurrence Matrix

- Hypothesis: Similar words tend to occur together and will have similar context.
- Example
 - Apple is a fruit. Mango is a fruit.
- Co-occurrence
 - For a given corpus, the co-occurrence of a pair of words is the number of times they have appeared together in a context window
- Context window
 - Context window is specified by a number and the direction



<https://www.analyticsvidhya.com/blog/2017/06/word-embeddings-count-word2veec/>

Co-occurrence Matrix

Corpus = He is not lazy. He is intelligent. He is smart.

| | He | is | not | lazy | intelligent | smart |
|-------------|----|----|-----|------|-------------|-------|
| He | 0 | 4 | 2 | 1 | 2 | 1 |
| is | 4 | 0 | 1 | 2 | 2 | 1 |
| not | 2 | 1 | 0 | 1 | 0 | 0 |
| lazy | 1 | 2 | 1 | 0 | 0 | 0 |
| intelligent | 2 | 2 | 0 | 0 | 0 | 0 |
| smart | 1 | 1 | 0 | 0 | 0 | 0 |

Context windows = 2

| | | | | | | | | | |
|----|----|-----|------|----|----|-------------|----|----|-------|
| He | is | not | lazy | He | is | intelligent | He | is | smart |
| He | is | not | lazy | He | is | intelligent | He | is | smart |
| He | is | not | lazy | He | is | intelligent | He | is | smart |
| He | is | not | lazy | He | is | intelligent | He | is | smart |
| He | is | not | lazy | He | is | intelligent | He | is | smart |

<https://www.analyticsvidhya.com/blog/2017/06/word-embeddings-count-word2vec/>

Co-occurrence Matrix

Matrix size

- $V \times V$
 - Not practical
- $V \times N$
 - N is a subset of V and can be obtained by removing irrelevant words like stop words
- $V \times k$
 - k is k principal components out of V using PCA

PCA to decompose Co-occurrence matrix

$X = U \cdot S \cdot V^T$ (Singular value decomposition)

U and S represent word vector

V presents word context

U is principal component.

$$\begin{array}{c} \hat{X} \\ \begin{pmatrix} x_{11} & x_{12} & \dots & x_{1n} \\ x_{21} & x_{22} & \dots & \\ \vdots & \vdots & \ddots & \\ x_{m1} & & & x_{mn} \end{pmatrix} \\ m \times n \end{array} \approx \begin{array}{c} U \\ \begin{pmatrix} u_{11} & \dots & u_{1r} \\ \vdots & \ddots & \\ u_{m1} & & u_{mr} \end{pmatrix} \\ m \times r \end{array} \begin{array}{c} S \\ \begin{pmatrix} s_{11} & 0 & \dots \\ 0 & \ddots & \\ \vdots & & s_{rr} \end{pmatrix} \\ r \times r \end{array} \begin{array}{c} V^T \\ \begin{pmatrix} v_{11} & \dots & v_{1n} \\ \vdots & \ddots & \\ v_{r1} & & v_{rn} \end{pmatrix} \\ r \times n \end{array}$$

$m \times m$
 $m \times n$
 $n \times n$

<https://www.analyticsvidhya.com/blog/2017/06/word-embeddings-count-word2vec/>

Frequency based Embedding

- Count Vector
 - Frequency : Number of times a word has appeared in the document
 - Presence : Count if the word appeared in the document
- TF-IDF Vector
- Co-occurrence Vector

<https://www.analyticsvidhya.com/blog/2017/06/word-embeddings-count-word2veec/>

Analogies

| | Man | Woman | King | Queen | Apple | Orange |
|--------|------|-------|-------|-------|-------|--------|
| Gender | -1 | 1 | -0.95 | 0.97 | 0 | 0.01 |
| Royal | 0.01 | 0.02 | 0.93 | 0.95 | -0.01 | 0 |
| Age | 0.03 | 0.02 | 0.7 | 0.69 | 0.03 | -0.02 |
| Food | 0.04 | 0.01 | 0.02 | 0.01 | 0.95 | 0.97 |

Man \rightarrow Woman vs. King \rightarrow ? (Queen)

$$\mathbf{e}_{\text{Man}} - \mathbf{e}_{\text{Woman}} \approx \mathbf{e}_{\text{King}} - \mathbf{e}_{\text{Queen}}$$

Linguistic regularities in continuous space word representations, Mikolov et. al., 2013,

Coursera: Deep learning Specialization, Andrew Ng

Analogies using word vectors

Man \rightarrow Woman vs. King \rightarrow ? (Queen)

$$\mathbf{e}_{\text{Man}} - \mathbf{e}_{\text{Woman}} \approx \mathbf{e}_{\text{King}} - \mathbf{e}_{\text{Queen}}$$

$$\mathbf{e}_{\text{Man}} - \mathbf{e}_{\text{Woman}} = \begin{bmatrix} -2 \\ 0 \\ 0 \\ 0 \end{bmatrix} \quad \mathbf{e}_{\text{King}} - \mathbf{e}_{\text{Queen}} = \begin{bmatrix} -2 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

Man:Woman as Boy:Girl
Ottawa:Canada as Nairobi:Kenya
Big:Bigger as Tall:Taller
Yen:Japan as Ruble:Russia

$$\mathbf{e}_{\text{Man}} - \mathbf{e}_{\text{Woman}} \approx \mathbf{e}_{\text{King}} - \mathbf{e}_w$$

Find word w : $\arg \max (\mathbf{e}_w, \mathbf{e}_{\text{King}} - \mathbf{e}_{\text{Man}} + \mathbf{e}_{\text{Woman}})$


Distance between two word vectors

(non-scaled) Euclidean distance

Defined in terms of inner product

$$\text{distance}(\mathbf{x}_i, \mathbf{x}_q) = \sqrt{(\mathbf{x}_i - \mathbf{x}_q)^T (\mathbf{x}_i - \mathbf{x}_q)}$$

$$= \sqrt{(x_i[1] - x_q[1])^2 + \dots + (x_i[d] - x_q[d])^2}$$



Cosine similarity – normalize

Similarity = $\frac{\sum_{j=1}^d x_i[j] x_q[j]}{\sqrt{\sum_{j=1}^d (x_i[j])^2} \sqrt{\sum_{j=1}^d (x_q[j])^2}}$

$\mathbf{x}_i^T \mathbf{x}_q = \cos(\theta)$


$\frac{\mathbf{x}_i^T \mathbf{x}_q}{\|\mathbf{x}_i\| \|\mathbf{x}_q\|} = \cos(\theta)$ (First normalize)

$\mathbf{a}^T \mathbf{b} = \|\mathbf{a}\| \|\mathbf{b}\| \cos(\theta)$

Diagram illustrating cosine similarity. Two vectors are shown in a 2D space defined by Feature 1 and Feature 2 axes. The angle between the vectors is labeled θ . The formula $\mathbf{a}^T \mathbf{b} = \|\mathbf{a}\| \|\mathbf{b}\| \cos(\theta)$ is shown.


Coursera: Machine Learning, Emily Fox & Carlos Guestrin

Pull out word vector from Embedding matrix



Coursera: Deep learning Specialization, Andrew Ng

Language model using word vector




Coursera: Deep learning Specialization, Andrew Ng

Language model using word vector

A neural probabilistic language model, Bengio et. al., 2003

$$\hat{P}(w_1^T) = \prod_{t=1}^T \hat{P}(w_t | w_1^{t-1})$$

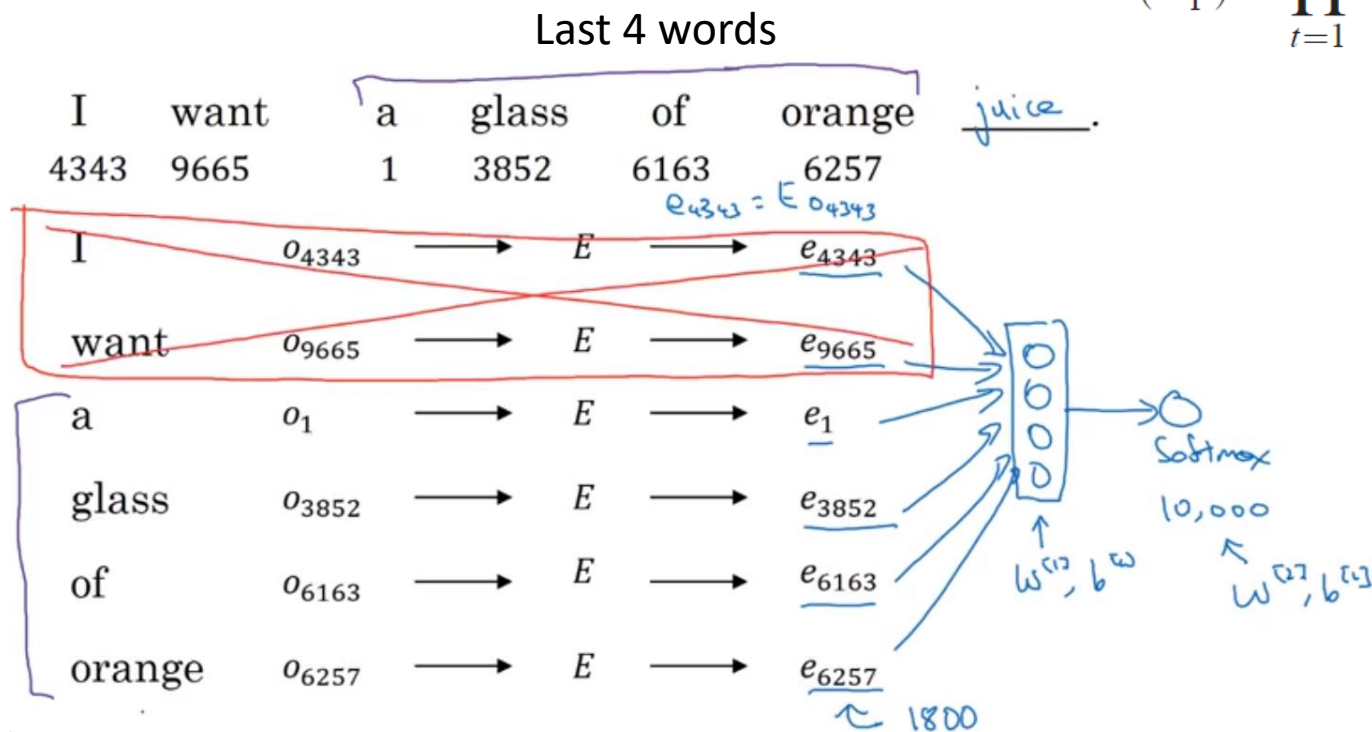


Coursera: Deep learning Specialization, Andrew Ng

Language model using word vector

A neural probabilistic language model, Bengio et. al., 2003

$$\hat{P}(w_1^T) = \prod_{t=1}^T \hat{P}(w_t | w_1^{t-1})$$



Coursera: Deep learning Specialization, Andrew Ng


Word2vec : Prediction based on Embedding

- CBOW (Continuous Bag of words)
 - Predicts the current word based on the context.
- Skip-gram
 - Predicts surrounding words given the current word.

Efficient Estimation of Word Representations in Vector Space, Tomas Mikolov et al., 2013

CBOW(Continuous Bag of words)


Predicts the current word based on the context



<https://ronxin.github.io/wevi/>

word vector :


the weight between the hidden layer and the output layer




Linear activation (no activation function between any layers)

<https://www.analyticsvidhya.com/blog/2017/06/word-embeddings-count-word2vec/>

Word2vec : Prediction based on Embedding



CBOW




Skip-gram

Efficient Estimation of Word Representations in Vector Space, Tomas Mikolov et al., 2013

Skip-gram

Predicts surrounding words given the current word



<https://www.analyticsvidhya.com/blog/2017/06/word-embeddings-count-word2vec/>

Context/target pair

I want a glass of orange juice to go along with my cereal.
target

Context:

Last 4 words : a glass of orange

4 words on left & right : a glass of orange, to go along with

Last 1 word : orange

Nearby 1 word → skip-gram

Skip-gram

I want a glass of orange juice to go along with my cereal.

Content

Orange

Orange

Orange

Target

Juice

Glass

to

Next 1 word

Left/right 2 windows

Skip-gram


How to sample content word?

- Randomly select → there is high chance to pick up most of stop-words (a, the, and, etc)
- Select from less-frequent words

Coursera: Deep learning Specialization, Andrew Ng

Softmax-classification

Vocab size = 10,000k



Θ_t : parameter associated with target t

Cost:
$$\mathcal{L}(\hat{y}, y) = - \sum_{i=1}^{10,000} y_i \log \hat{y}_i$$

Coursera: Deep learning Specialization, Andrew Ng

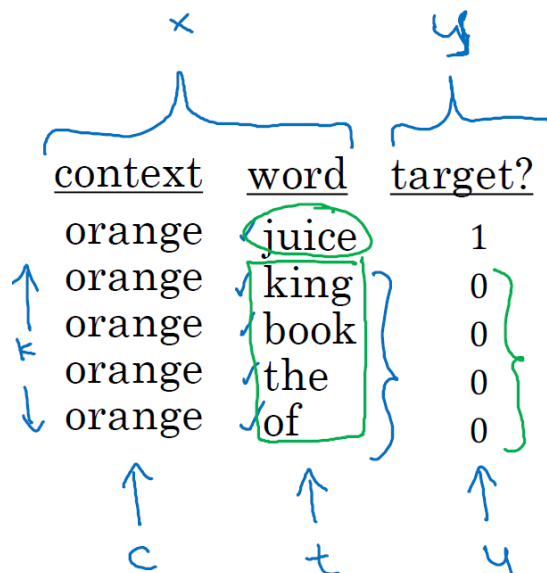
Negative sampling

I want a glass of orange juice to go along with my cereal.

Binary classification

$$P(y=1 | c, t) = \sigma(\theta_c^T e_c)$$

Sigmoid function



Randomly pick up
from dictionary

or

$$P(w_i) = \frac{f(w_i)^{3/4}}{\sum_{j=1}^{10,000} f(w_j)^{3/4}}$$

Coursera: Deep learning Specialization, Andrew Ng

Word2vec

vocabulary_size=7 and embedding_size=3

| | | | |
|-------------------|------|------|------|
| <i>anarchism</i> | 0.5 | 0.1 | -0.1 |
| <i>originated</i> | -0.5 | 0.3 | 0.9 |
| <i>as</i> | 0.3 | -0.5 | -0.3 |
| <i>a</i> | 0.7 | 0.2 | -0.3 |
| <i>term</i> | 0.8 | 0.1 | -0.1 |
| <i>of</i> | 0.4 | -0.6 | -0.1 |
| <i>abuse</i> | 0.7 | 0.1 | -0.4 |

| | Man | Woman | King | Queen | Apple | Orange |
|--------|------|-------|-------|-------|-------|--------|
| Gender | -1 | 1 | -0.95 | 0.97 | 0 | 0.01 |
| Royal | 0.01 | 0.02 | 0.93 | 0.95 | -0.01 | 0 |
| Age | 0.03 | 0.02 | 0.7 | 0.69 | 0.03 | -0.02 |
| Food | 0.04 | 0.01 | 0.02 | 0.01 | 0.95 | 0.97 |

<http://adventuresinmachinelearning.com/word2vec-tutorial-tensorflow/>

GloVe: Global Vectors for Word Representation

GloVe: Global Vectors for Word Representation, Jeffrey Pennington et al., 2014

I want a glass of orange juice to go along with my cereal.



Shallow Window-based methods (making predictions within local context windows)

GloVe: Model

GloVe: Global Vectors for Word Representation, Jeffrey Pennington et al., 2014

Loss function
$$J = \sum_{i,j=1}^V f(X_{ij}) (w_i^T \tilde{w}_j + b_i + \tilde{b}_j - \log X_{ij})^2$$

minimize
$$\sum_{i=1}^{10,000} \sum_{j=1}^{10,000} f(x_{ij}) (\Theta_i^T e_j + b_i + b_j' - \log x_{ij})^2$$

1. $f(0) = 0$. If f is viewed as a continuous function, it should vanish as $x \rightarrow 0$ fast enough that the $\lim_{x \rightarrow 0} f(x) \log^2 x$ is finite.
2. $f(x)$ should be non-decreasing so that rare co-occurrences are not overweighted.
3. $f(x)$ should be relatively small for large values of x , so that frequent co-occurrences are not overweighted.




Figure 1: Weighting function f with $\alpha = 3/4$.

GloVe: Global Vectors for Word Representation

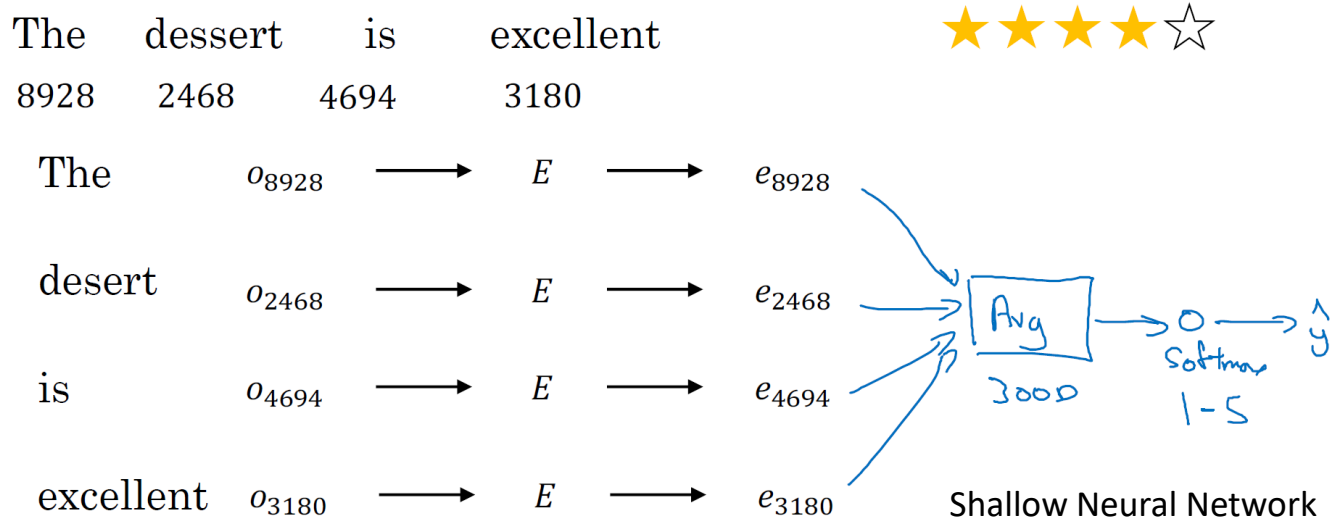
GloVe: Global Vectors for Word Representation, Jeffrey Pennington et al., 2014

- Unsupervised word representation
- Combined
 - Count-based method
 - Prediction-based method
- Outperforms
 - Word analogies
 - Word similarity
 - Named entity recognition

Sentiment classification problem

| x | y |
|--|-------|
| The dessert is excellent. | ★★★★☆ |
| Service was quite slow. | ★★☆☆☆ |
| Good for a quick meal, but nothing special. | ★★★☆☆ |
| Completely lacking in good taste, good service, and good ambience. | ★☆☆☆☆ |

Simple sentiment classification model

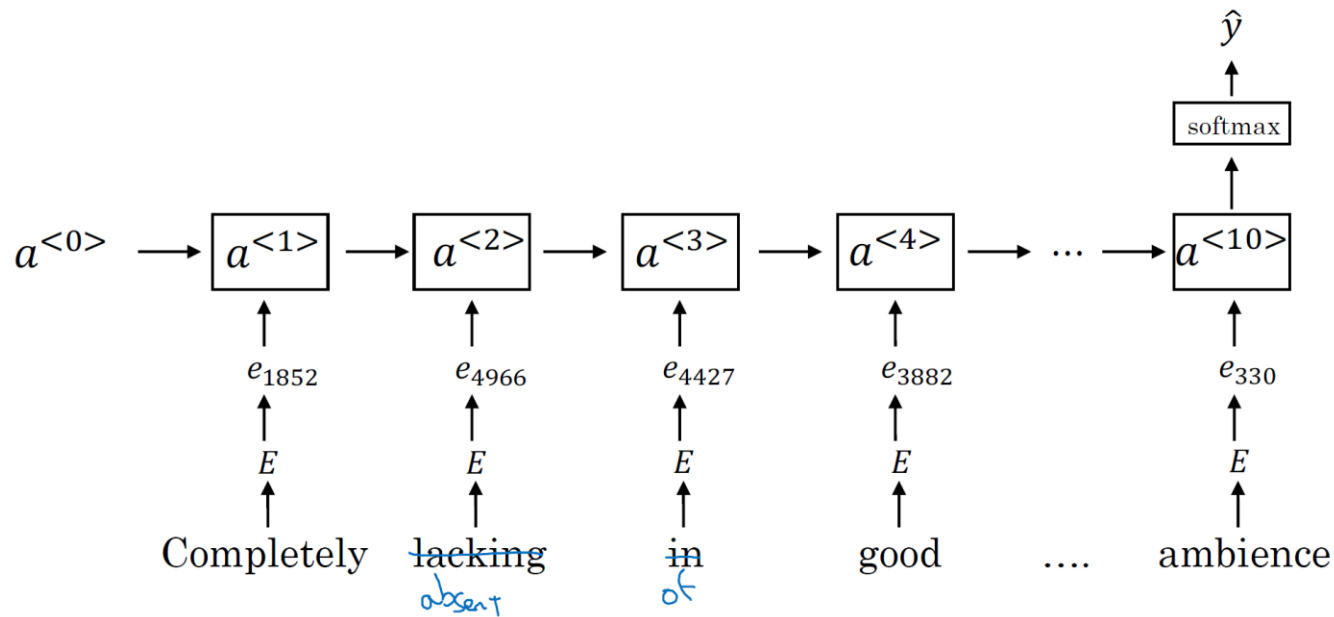


Problem in Simple sentiment classification model

Completely lacking in good taste, good service, and good ambience.



RNN for sentiment classification



Summary

- Character-based token → 1D convolution(word) + pooling(n-gram) → Neural Network → softmax
- Count vector or TF-IDF → Neural Network → softmax
- Word vector (word2vec, glove) → Neural Network → softmax
- Word-embedding (word2vec, glove) → RNN